

SAFETYANALYST: INTERPRETABLE, TRANSPARENT, AND STEERABLE LLM SAFETY MODERATION

ABSTRACT

The ideal LLM content moderation system would be both structurally interpretable (so its decisions can be explained to users) and steerable (to reflect a community’s values or align to safety standards). However, current systems fall short on both of these dimensions. To address this gap, we present SAFETYANALYST, a novel LLM safety moderation framework. Given a prompt, SAFETYANALYST creates a structured “harm-benefit tree,” which identifies 1) the actions that could be taken if a compliant response were provided, 2) the harmful and beneficial effects of those actions (along with their likelihood, severity, and immediacy), and 3) the stakeholders that would be impacted by those effects. It then aggregates this structured representation into a harmfulness score based on a parameterized set of safety preferences, which can be transparently aligned to particular values. **To demonstrate the power of this framework, we develop, test, and release a prototype system, SAFETYREPORTER, including a pair of LMs specializing in generating harm-benefit trees through symbolic knowledge distillation and an interpretable algorithm that aggregates the harm-benefit trees into safety labels. SAFETYREPORTER is trained on 18.5 million harm-benefit features generated by SOTA LLMs on 19k prompts.** On a comprehensive set of prompt safety benchmarks, we show that our system (average F1=0.75) outperforms existing LLM safety moderation systems (average F1<0.72) on prompt safety classification, while offering the additional advantages of interpretability and steerability.

1 INTRODUCTION

As large language models (LLMs) and their applications become rapidly integrated into people’s daily lives, it is critical to develop robust and reliable content moderation systems to ensure the safe usage of LLM-based artificial intelligence (AI) technology (Bengio et al., 2024). Recently, Dalrymple et al. (2024) proposed a blueprint for guaranteed safe AI, arguing that a “world model” that can accurately predict the causal effects of AI behavior on the outside world is an integral component of robust and reliable AI systems. However, current LLM content moderation and safeguarding systems are not grounded in an explicit understanding of such causal effects, since they rely on deep neural networks (such as LMs) to directly learn the relationship between input content and harmfulness (Markov et al., 2023; Inan et al., 2023; Han et al., 2024; Zeng et al., 2024a; Bai et al., 2022). Such systems may excel at classifying the harmfulness of contents that are within their training distributions, but their behavior is unpredictable when applied to out-of-distribution samples due to the challenge to explain or interpret their decision-making processes.

Moreover, as AI technology reaches diverse human populations (e.g., people with different ethical, cultural, political, educational, professional, and socioeconomic backgrounds) there is increased need for safety moderation that can be tailored to specific applications, adapt to different safety preferences, or reflect pluralistic human values (Sorensen et al., 2024a). AI safety criteria change for different user demographics. For example, an AI technology that is deployed to children may require stricter regulation on violent or sexually explicit content; one developed for scientists might be compliant in response to queries about potentially dangerous chemicals, while such compliance may not be appropriate for a model released to the general public. Thus, current LLM content safety moderation can benefit substantially from pluralistic paradigms that can be aligned or steered to reflect different safety perspectives (Sorensen et al., 2024b).

054
055
056
057
058
059
060
061
062
063
064
065
066
067
068
069
070
071
072
073
074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
102
103
104
105
106
107




Figure 1: Overview of the **SAFETYREPORTER** system that implements the **SAFETYANALYST** framework on the **prompt safety classification** task. We generated extensive harm-benefit feature data using SOTA LLMs (GPT-4o, Gemini-1.5-Pro, Llama-3.1-70B-Instruct, Llama-3.1-405B-Turbo, and Claude-3.5-Sonnet) on 19k user prompts through chain-of-thought prompting. We embedded each prompt in a hypothetical AI language model usage scenario and instructed the LLMs to enumerate all stakeholders who may be impacted, any potentially harmful/beneficial actions that may impact the stakeholders, and the effects each action may cause to each stakeholder. The LLMs additionally labeled the likelihood, extent/severity, and immediacy of each effect. These harm-benefit features were then used to train two specialist models — one to generate harms and one to generate benefits (together part of **SAFETYREPORTER**) — through symbolic knowledge distillation via supervised fine-tuning of Llama-3.1-8B-Instruct. Given any prompt, **SAFETYREPORTER** efficiently generates an interpretable harm-benefit tree. The harms and benefits are weighted and traded off by an aggregation algorithm to calculate a harmfulness score, which can be directly translated into content safety labels or refusal decisions. Steerability can be achieved by aligning the weights in the aggregation algorithm to a user’s or community’s preference or to principled safety standards.

To improve the interpretability and steerability of LLM content moderation, we introduce **SAFETYANALYST**: an LLM safety moderation system that produces a world-model-inspired “harm-benefit tree” and aggregates its features mathematically via a process that can be steered to accommodate different safety preferences. While existing AI safety **content moderation tools rely on opaque systems which categorize prompts as harmful without fully interpretable further explanation** (Zeng et al., 2024b; Xie et al., 2024; Han et al., 2024; Ji et al., 2024; Mazeika et al., 2024), **SAFETYANALYST** is grounded in the fundamental principles of cost-benefit analysis (Arrow et al., 1996), explicitly representing what *actions* may cause which harmful or beneficial *effects* for different *stakeholders*. Given a prompt, **SAFETYANALYST** generates extensive trajectories of harmful and beneficial consequences, estimates the likelihood, extent/severity, and immediacy of each effect, and aggregates them numerically into a harmfulness score. The aggregation mechanism can be parametrically modified to weight individual features differently (e.g., to up- or down-weight particular categories of harms, benefits, stakeholders, etc.). Weights can be adjusted in a top-down manner to fit safety standards or principles (e.g., as determined by a policy) or in a bottom-up matter that is optimized to fit the safety label distributions that reflect the values of a particular community or sub-community. Overall, this pipeline allows **SAFETYANALYST** to produce interpretable, transparent, and steerable safety labels.

We implemented the conceptual **SAFETYANALYST** framework into a system for prompt harmfulness classification, named **SAFETYREPORTER**. Using 19k harm-benefit trees generated by a mixture of

state-of-the-art (SOTA) LLMs containing 18.5 million features, we fine-tuned an open-weight LM to specialize in generating harm-benefit features. To perform prompt classification, we optimized the parameters of our mathematical aggregation algorithm to the harmful and benign prompt labels provided by WildJailbreak, a large-scale prompt dataset containing synthetic benign and harmful prompts generated based on 13 risk categories (Jiang et al., 2024). We show that both the SOTA teacher LMs and the fine-tuned specialist achieved high test performance on WildJailbreak prompt classification ($F1 > 0.84$, $AUPRC > 0.89$, and $AUROC > 0.88$). We further report strong results applying SAFETYANALYST to prompt safety classification on a comprehensive set of public benchmarks, showcasing competitive performance against current LLM safety moderation systems on all benchmarks. On average, our system ($F1 = 0.75$) outperformed existing counterparts ($F1 < 0.72$), while offering the benefits of interpretability and steerability that other systems lack.

Contributions. In this paper, we introduce SAFETYANALYST, a novel conceptual framework for LLM safety content moderation that offers more interpretability, transparency, and steerability than existing approaches. The framework proposes a method to surface structured harmful and beneficial effects of a user prompt (in the form of “harm-benefit trees”), which can then be mathematically aggregated according to their weights. To facilitate use of this framework, we train and release SAFETYREPORTER, an open-source pair of LMs that specialize in the task of harm-benefit tree creation, which we evaluate against SOTA content-moderation tools showing competitive performance. In addition, we release a series of other artifacts that enable researchers and engineers to build on SAFETYANALYST: a large-scale dataset of 18.5 million safety features (organized in as harm-benefit trees) generated by SOTA LLMs on 19k prompts, the first taxonomies of harmful and beneficial effects for AI safety, and a feature aggregation algorithm that can be steered to align with a given safety content label distribution or with top-down safety standards.

2 THE SAFETYANALYST FRAMEWORK AND SAFETYREPORTER SYSTEM

SAFETYANALYST breaks down the problem of content classification into sub-tasks (Figure 1). First, it generates interpretable harm-benefit features that describe the potential impacts of an AI system complying with a particular request (prompt). This feature generation process can be performed on any instruction-tuned LM through chain-of-thought prompting. Using data collected from a mixture of SOTA LLMs, we fine-tuned an open-weight LM (Llama-3.1-8B-Instruct) to specialize in efficient feature generation. Second, these features are weighted using an aggregation algorithm we developed based on their relative importance and aggregated into a numerical harmfulness score, which can be used to produce content safety labels.

2.1 HARM-BENEFIT FEATURE GENERATION

Given a prompt and a scenario where the AI language model complies with the user request, an LM extensively generates features (Figure 2) including all stakeholders (individuals, groups, communities, and entities in society that may be affected), harmful and beneficial actions that may impact each stakeholder, harmful and beneficial effects that may be caused by each action on each stakeholder, and the likelihood (low, medium, or high), extent/severity (minor, significant, substantial, or major), and immediacy (immediate or downstream) of each effect. Harmful actions are generated in accordance with (and classified by) the AIR 2024 risk taxonomy (Zeng et al., 2024b), an extensive categorization of harmful actions that could result from interaction with an LM, derived from worldwide governmental and corporate policies. Beneficial actions are generated in free text. Due to the lack of formal characterization of harmful and beneficial *effects* in the AI safety literature, we defined a novel hierarchical taxonomy, drawing on the theories of basic/primary goods of two influential contemporary moral philosophers: Bernard Gert (Gert, 2004) and John Rawls (Rawls, 2001). See Appendix A for complete taxonomies.

We used a diverse mixture of SOTA LLMs including GPT-4o (Achiam et al., 2023), Gemini-1.5-Pro (Team et al., 2023), Llama-3.1-70B-Instruct, Llama-3.1-405B-Instruct-Turbo (Dubey et al., 2024), and Claude-3.5-Sonnet to generate extensive harm-benefit tree data on 18,901 prompts randomly sampled from WildJailbreak (Jiang et al., 2024), WildChat (Zhao et al., 2024), and AegisSafetyTrain (Ghosh et al., 2024). Table 3 in Appendix B shows the breakdown of prompt distribution over the




Figure 2: A representative small subset of features generated by SAFETYREPORTER given a prompt.

datasets for all LLMs. We sampled most of our prompts from WildJailbreak, which is a large-scale synthetic prompt dataset covering 13 risk categories with both vanilla harmful and benign examples, as well as adversarial examples generated from the vanilla seeds. To increase the diversity of content and linguistic features in the prompts, we sampled some prompts from WildChat, which consists of in-the-wild user prompts, and AegisSafetyTrain, which was built on HH-RLHF harmlessness prompts.

Overall, the LLMs generated rich harm-benefit features that follow a tree-like structure: more than 10 stakeholders per prompt, 3-10 actions per stakeholder, 3-7 effects per action, varying between models and prompt classes in WildJailbreak (Table 4 in Appendix B). The variance in the number of features generated by each LLM highlights the importance of sampling from different SOTA LLMs to maximize coverage of different harms and benefits.

2.2 SAFETYREPORTER: AN OPEN-SOURCE PAIR OF SPECIALIST MODELS FOR HARM AND BENEFIT FEATURE GENERATION

To enable fast, cheap, and high quality harm-benefit feature generation, we trained an open-weight LM (Llama-3.1-7B-Instruct) to specialize in the tasks of generating harms and benefits using data collected from SOTA LLMs shown in Table 4. We applied supervised fine-tuning using qlora (Dettmers et al., 2024) to distill the knowledge about harmful and beneficial features of our interest from the teacher models (SOTA LLMs) into the student model (West et al., 2021). We trained one specialist model to generate harm-trees and another for benefit-trees, which can be combined into the full harm-benefit tree structure (Figure 2). Due to the extensive combined lengths of our taxonomies and the harm-benefit trees generated by teacher LLMs, we fine-tuned two specialists instead of one so that the inputs and outputs could jointly fit into the context window defined by our hardware constraints (context window length of 18,000 tokens on 8 NVIDIA H100 GPUs). The two student models that specialize in harm and benefit feature generation are integral components of SAFETYREPORTER.

Table 1: Performance of models operating in the SAFETYANALYST framework on the WildJailbreak prompt safety classification task. Three “teacher” models as well as SAFETYREPORTER, the student, were tested. Each model generated a harm-benefit tree for each prompt, which was then passed to the model-specific aggregation algorithm, which was used to generate a prompt classification.

Metric	GPT-4o	Gemini-1.5-Pro	Llama-3.1-70B	SAFETYREPORTER
F1	91.8	87.7	88.1	84.7
AUPRC	91.7	92.0	96.6	89.0
AUROC	94.7	92.5	95.9	88.4

We trained SAFETYREPORTER on all data generated by the teacher models shown in Table 4 except that we randomly down-sampled the WildJailbreak data from Llama-70B to 1,000 vanilla harmful and 1,000 vanilla benign prompts. Additionally, to increase the robustness of SAFETYREPORTER to adversarial attacks (e.g., jailbreaks), we augmented the training dataset with adversarial prompts from WildJailbreak, which contains synthetic adversarial prompts created based on the vanilla prompts using in-the-wild jailbreak techniques. We randomly sampled 6,368 adversarial prompts that corresponded to the vanilla prompts (at most one adversarial prompt per vanilla prompt) used in data generation, and augmented the training dataset by pairing them with the harm-benefit trees of the corresponding vanilla prompts.

To evaluate the quality of generated harm-benefit features, we collected human annotation data from 126 prolific workers on their agreement with the generated stakeholders, harmful/beneficial effects, and the likelihoods, extents, and immediacies of the effects. Annotators showed broad agreement on the plausibility of the harm-benefit features (see Table 5 in Appendix C for results, Figure 4 in Appendix C for interface design, and Section 4 for further discussion).

2.3 MATHEMATICAL FEATURE AGGREGATION

We mathematically formalize a feature aggregation algorithm for quantifying the harmfulness (H) of a prompt over features generated by a SAFETYANALYST model parameterized by W and γ :

$$H(\text{prompt} \mid W, \gamma) = \sum_{\text{Stakeholder}} \sum_{\text{Action}} \sum_{\text{Effect}} W_{\text{Action}} \cdot W_{\text{Likelihood}} \cdot W_{\text{Extent}} \cdot W_{\text{Immediacy}},$$

where W is a set of weights for the 16 second-level action categories in the AIR 2024 taxonomy and relative importance weights of different extents and likelihoods. γ includes discount factors for downstream (vs. immediate) and beneficial (vs. harmful) effects. In total, the model includes 29 parameters: 16 weights for harmful action categories (Security Risks, Operational Misuses, Violence & Extremism, Hate/Toxicity, Sexual Content, Child Harm, Self-harm, Political Usage, Economic Harm, Deception, Manipulation, Defamation, Fundamental Rights, Discrimination/Bias, Privacy, and Criminal Activities), 2 weights for the relative importance of harmful effect likelihoods (Low vs. Medium and Medium vs. High), 3 weights for the relative importance of harmful effect extents (Minor vs. Significant, Significant vs. Substantial, and Substantial vs. Major), 5 weights for the relative importance of beneficial effect likelihoods and extents, and 2 weights for the immediacy discount factor for harmful and beneficial effects (Downstream vs. Immediate). By default, $W_{\text{High likelihood}} = 1$, $W_{\text{Major extent}} = 1$, and $W_{\text{Immediate}} = 1$ for all harms and $W_{\text{Beneficial action}} = -1$.

2.4 FEATURE WEIGHT ALIGNMENT

To translate the numerical harmfulness score H computed over features generated by some SAFETYANALYST model into a safety label for prompt classification, we aligned the aggregation algorithm to the ground-truth labels from the WildJailbreak dataset on harm-benefit trees generated by teacher and student models by optimizing W and γ within $[0, 1]$ using maximum-likelihood estimation over the analytical likelihood of $\sigma(H)$. This procedure optimized the weights to minimize the discrepancy between true and predicted safety labels. At inference time, the weights were frozen at their optimal values. Table 1 shows the classification performance (measured by the F1 score,

AUPRC, and AUROC and presented in percentage) of different teacher and student SAFETYANALYST models (GPT-4o, Gemini-1.5-Pro, Llama-3.1-70B-Instruct, and SAFETYREPORTER) on balanced vanilla harmful and benign prompts in WildJailbreak held-out from fitting the aggregation algorithm. All models achieved high classification performance, with the lowest F1 = 84.7, AUPRC = 89.0, and AUROC = 88.4. Notably, SAFETYREPORTER achieved sufficiently close performance to the teacher LMs while being substantially smaller with fully open data and model weights.

The optimized parameter values are illustrated in Figure 3. Among the harmful actions summarized by level-2 risk categories in the AIR 2024 taxonomy (Zeng et al., 2024b), Self-harm weighted the highest, followed by Criminal Activities and Political Usage. High likelihood, immediate effects dominated the aggregation, with near-zero weights for medium and low likelihood or downstream effects, except for medium likelihood harmful effects. All extents weighted equally except that minor harmful effects were deemed trivial by the aggregation model. Overall, aggregation was driven by harmful effects, as evident by the low relative importance of a beneficial effect compared to a harmful effect (13.4%).

3 APPLYING SAFETYREPORTER TO PROMPT SAFETY CLASSIFICATION

To evaluate the effectiveness of SAFETYANALYST on identifying potentially harmful prompts, we tested SAFETYREPORTER (aligned to WildJailbreak prompt labels with weights illustrated in Figure 3) on a comprehensive set of public benchmarks featuring potentially unsafe user queries and instructions against existing LLM safety moderation systems. Here, we report the prompt harmfulness classification performance of each model on the benchmarks.

3.1 EVALUATION SETUP

Benchmarks. We tested SAFETYREPORTER and relevant baselines on 6 publicly available prompt safety benchmarks, including SimpleSafetyTests (100 prompts; Vidgen et al. 2023), HarmBench-Prompt standard test set (159 prompts; Mazeika et al. 2024), WildGuardTest (960 vanilla and 796 adversarial prompts; Han et al. 2024), AIR-Bench-2024 (5,694 prompts; Zeng et al. 2024c), and SORRY-Bench (9,450 prompts; Xie et al. 2024). These benchmarks represent a diverse and comprehensive selection of unsafe prompts, including manually crafted prompts on highly sensitive and harmful topics (SimpleSafetyTests), standard behavior that may elicit harmful LLM responses (HarmBench), adversarial prompts (WildGuardTest), benign prompts (WildGuardTest), prompts that may challenge government regulations and company policies (AIR-Bench-2024), and unsafe prompts that cover granular risk topics and linguistic characteristics (SORRY-Bench). Since our system focuses on identifying prompts that would be unsafe to respond to, rather than the harmfulness in the prompt content per se, we did not include benchmarks in which prompts were labeled




Figure 3: Optimized SAFETYREPORTER aggregation feature weights, fitted to balanced WildJailbreak prompt labels. Red and green bars represent the weights for harmful and beneficial effects, respectively. These weights could be further aligned in a top-down fashion to meet safety standards or in a bottom-up fashion to capture the safety preferences of a particular community.

for the latter, such as the OpenAI moderation dataset (Markov et al., 2023), ToxicChat (Lin et al., 2023), and AegisSafetyTest (Ghosh et al., 2024).

Baselines. We compare SAFETYREPORTER to 9 existing LLM safety moderation systems: OpenAI moderation endpoint (Markov et al., 2023), LlamaGuard, LlamaGuard-2, LlamaGuard-3 (Inan et al., 2023), Aegis-Guard-Defensive, Aegis-Guard-Permissive (Ghosh et al., 2024), ShieldGemma-2B, ShieldGemma-9B, ShieldGemma-27B (Zeng et al., 2024a), and WildGuard (Han et al., 2024). Additionally, we report zero-shot GPT-4 performance (Achiam et al., 2023). In Appendix D, we provide detailed descriptions of all baselines evaluated.

We referenced Han et al. (2024)’s evaluation results where applicable and additionally tested models and benchmarks that they did not feature with temperature set to 0. We were unable to fairly evaluate Llama-Guard, Aegis-Guard-Defensive, and Aegis-Guard-Permissive (both Aegis-Guards are tuned Llama-Guard models) on SORRY-Bench, since the lengths of 457 prompts in SORRY-Bench exceeded the Llama-2 context window limit of 4,096 tokens (Touvron et al., 2023). For each model, we computed an average F1 score across benchmarks weighted by the number of prompts in each benchmark dataset. Experiments using open-weight models were run on one NVIDIA H100 GPU with batched inference using vllm (Kwon et al., 2023).

3.2 EVALUATION RESULTS

SAFETYREPORTER outperforms existing LLM safety moderation systems on prompt harmfulness classification. Table 2 shows our evaluation results, measured by the F1 score (denoted in percentage). SAFETYREPORTER achieved competitive performance on all benchmarks compared to existing LLM safety moderation systems, with the highest overall F1 score of 75.4, exceeding the second highest score of 71.7 by WildGuard. Notably, SAFETYREPORTER’s performance was zero-shot, since it was not trained on or aligned to any training datasets of the benchmarks, whereas WildGuard was trained on the WildGuardTrain set. Nonetheless, GPT-4’s classification performance was better than all the LLM moderation models with an F1 score of 81.6. In Appendix D.3, we show that GPT-4’s outstanding performance on SORRY-Bench was driven by its better capability to identify potentially unsafe prompts encoded or encrypted in Atbash and Caesar ciphers. SAFETYREPORTER outperformed other baselines on identifying potentially unsafe prompts against Persuasion Techniques (Authority Endorsement, Evidence-based Persuasion, Expert Endorsement, Logical Appeal, and Misrepresentation).

Table 2: F1 scores of prompt harmfulness classification on public benchmarks. The average was computed over all benchmarks weighted by the number of examples in each dataset. The highest average score is emphasized in bold and the second highest underlined.

Model	SimpS-Tests	Harm-Bench	WildGuardTest		AIR-Bench	SORRY-Bench	Average
			Vani.	Adv.			
OpenAI Mod. API	63.0	47.9	16.3	6.8	46.5	42.9	41.1
Llama-Guard	93.0	85.6	70.5	32.6	44.7	-	-
Llama-Guard-2	95.8	91.8	85.6	46.1	74.9	53.9	62.9
Llama-Guard-3	99.5	98.4	86.7	61.6	68.8	59.1	64.6
Aegis-Guard-D	100	93.6	82.0	74.5	83.4	-	-
Aegis-Guard-P	99.0	87.6	77.9	62.9	62.5	-	-
ShieldGemma-2B	99.5	100	62.2	59.2	28.6	18.5	27.4
ShieldGemma-9B	83.7	77.2	61.3	35.8	28.6	39.0	37.3
ShieldGemma-27B	85.7	74.8	62.4	43.0	32.0	42.3	40.6
WildGuard	99.5	99.7	91.7	85.5	87.6	58.2	71.7
GPT-4	100	100	93.4	81.6	84.5	78.2	81.6
SAFETYREPORTER	95.2	94.4	88.3	73.7	83.0	69.1	<u>75.4</u>

Inference-time compute. Due to the extensiveness of the harm-benefit trees generated by SAFETYREPORTER for each prompt (Figure 2; Table 4), it requires more inference-time compute than other baselines that only produce safety labels. On the same computing infrastructure, SAFETYREPORTER averaged 6.12 seconds per prompt and WildGuard 0.22 second per prompt. Therefore, the current instantiation of the SAFETYANALYST framework (i.e., as implemented in SAFETYREPORTER) is best reserved for cases where steerable and interpretable safety moderation is highly valued over compute usage at inference time. Future work should explore how other implementations of the SAFETYANALYST framework on different architectures could reduce computational intensity. Moreover, SAFETYREPORTER’s own inference could be substantially accelerated by parallel computing. Finally, if a faster system were desired, a promising approach would be to selectively lesion the harm-benefit trees to only preserve the most helpful features. As a demonstration of this approach, we systematically ablated different dimensions of the harm-benefit trees and report the model’s performance on WildGuardTest and WildJailbreak (Appendix D.3). Our results show that harms contributed more than benefits, and likelihood more than extent and immediacy in the aggregation algorithm fitted to WildJailbreak. However, since this observation may not hold true for all datasets and tasks (particularly for those where disagreements among annotators are likely), we generated the full harm-benefit tree in the current work for generality.

3.3 ADDITIONAL BENEFITS OF SAFETYREPORTER

Interpretability. Although SAFETYREPORTER achieved outstanding performance on prompt safety classification, its most critical advantage is the interpretability of its decision-making process compared to black-box systems, including all the baselines in Table 2. This interpretability is two-folded: first, the features, on which the safety decisions are based solely, are explicitly generated by SAFETYREPORTER and semi-structured (i.e., on carefully curated dimensions, including stakeholder, harm, benefit, action, effect, extent, likelihood, and immediacy); second, these features are aggregated using a white-box algorithm with transparent mechanisms and interpretable feature weights that quantify the importance of corresponding feature values (Figure 3). Even though LLMs (such as GPT-4) can generate explanations for their decisions, there remains a lack of interpretability in *how* the decisions are reached and there is no reliable causal relationship between the explanation and the safety prediction. Our strong evaluation results in Table 2 suggest that our simple but interpretable features and aggregation mechanisms contain sufficient information for decision-making on content safety. Appendix E includes a detailed example of the full decision-making process of SAFETYREPORTER, highlighting its interpretability and transparency.

Steerability. In addition, SAFETYREPORTER’s aggregation algorithm is defined by a set of transparent, interpretable parameter weights. The weights of the parameters we report in Figure 3 reflect the values of the annotators who provided the labels for the WildJailbreak dataset, for which the algorithm was optimized. However, one central strength of the SAFETYANALYST approach is that the aggregation algorithm allows different safety features to be up- or down-weighted for top-down adjustments, or fitted to a customized safety label distribution for bottom-up adjustments (e.g., personalized safety alignment). Bottom-up adjustments of weights can be achieved by fitting the aggregation model to a safety label distribution produced by an individual or group; the resulting parameters would be aligned to the values expressed in the labels. We provide concrete explanations for how to operationalize top-down weight adjustments in the case study in Appendix E.

4 RELATED WORK

Existing LLM content moderation systems. While there are many ways to approach AI safety, SAFETYANALYST is designed to do so through *content moderation*, the goal of which is to ensure that an AI system “avoids unsafe, illegal outputs” (Huang et al., 2024). Existing LLM content moderation systems include WildGuard (Han et al., 2024), ShieldGemma (Zeng et al., 2024a), AegisGuard (Ghosh et al., 2024), LlamaGuard (Inan et al., 2023), and the OpenAI moderation endpoint (Markov et al., 2023). These systems are LM-based classifiers that can categorize content risk, including user prompts. Except for minor variations, each of these systems is structured similarly: a general-purpose LLM is trained on a large dataset that links user prompts to harmfulness labels.

The resulting content moderation systems then can classify prompts as harmful or not based on the training it received (see Appendix D.1 for details). Although some systems built in this way can achieve high classification accuracy on prompt safety benchmarks (e.g., classifying a prompt as harmful or benign), their internal decision mechanisms are challenging to interpret, which limits their reliability and generalizability. **There is no straight-forward way to determine why a prompt was classified as harmful by one of these systems.** Furthermore, due to the lack of modularity in their architectures, they cannot be easily steered to reflect different safety perspectives **beyond expensive and time-consuming re-training or fine-tuning processes.**

LLM content risk. Prior work has characterized LLM content safety based on the potential risk of the content, including the user input to the LLM, which may include jailbreak attacks, and the LLM output, on general and specific applications (Bai et al., 2022; Shen et al., 2023; Huang et al., 2024; Ji et al., 2024; Walker et al., 2024). The AI safety literature has relied on risk taxonomies to categorize unsafe content. Recent work has built on standard risk categories (Weidinger et al., 2022) to include more fine-grained categories (Wang et al., 2023; Tedeschi et al., 2024; Xie et al., 2024; Brahman et al., 2024), achieve comprehensive coverage (Vidgen et al., 2024), and incorporate government regulations and company policies (Zeng et al., 2024b). Our system relies on the taxonomy developed by Zeng et al. (2024b), selected for its comprehensive and fine-grained nature. Overall, these taxonomies describe the unsafe nature of a prompt or unsafe actions that might result from a prompt being answered. To our knowledge, no prior work exists that proposes formal taxonomies for the downstream *effects* of unsafe prompts (as opposed to *actions*; see Appendix A for our taxonomies of harmful and beneficial effects).

Symbolic knowledge distillation. We distilled a pair of small, expert LMs (SAFETYREPORTER) to create structured harm-benefit trees, the core of our interpretable framework. The symbolic knowledge distillation strategy leverages diffuse knowledge gained by large, generalist (and often proprietary) models to create a more compact expert student model that excels at one particular task (Xu et al., 2024; West et al., 2021; Tang et al., 2019). This strategy is useful (among other reasons) to generate rich, structured data that is too costly or labor-intensive for humans to do by hand (West et al., 2021). Indeed, prior work shows that symbolic knowledge distillation from machine teachers can exceed the quality of human-authored symbolic knowledge (West et al., 2021; Jung et al., 2023). Compared to the teacher models, our SAFETYREPORTER uses less time, memory, compute, and cost while achieving comparable performance, and it will be openly released for public use in LLM moderation contexts.

Pluralistic alignment for LLM safety. Although current LLM safety moderation systems are yet to be pluralistically aligned, recent interest in value pluralism Sorensen et al. (2024a) has given rise to rapid developments of pluralistic alignment approaches for LLMs. Lera-Leri et al. (2022) formalized an aggregation method for value systems inspired by the social choice literature. Feng et al. (2024) outlined a more general framework based on multi-LLM collaboration, in which an LLM can be aligned to specialized community LMs for different pluralism objectives. Other methods have been proposed for learning distributions of human preferences rather than the majority (Siththaranjan et al., 2023; Chen et al., 2024). **Additionally, some recent work has featured individualized human preference data, including the DICES dataset (Aroyo et al., 2024) and the PRISM alignment project (Kirk et al., 2024), paving the path to pluralistically or personally aligned LLM systems.**

5 CONCLUSION

We introduce SAFETYANALYST, a novel conceptual framework **based on LM-generated, semi-structured harm-benefit trees** for interpretable, transparent, and steerable LLM content safety moderation. We operationalized the pipeline of harm-benefit tree data generation **through chain-of-thought prompting**, symbolic knowledge distillation, and weighted feature aggregation to implement a system for prompt safety classification. Our system achieved SOTA performance on a comprehensive set of prompt safety benchmarks, promising strong potential in real-world LLM safety applications.

Our application of SAFETYANALYST and SAFETYREPORTER to a comprehensive set of prompt safety benchmarks shows SOTA performance compared to existing LLM safety moderation systems.

The current implementation of SAFETYANALYST focuses on prompt harmfulness classification, which can help an AI system determine if a user prompt should be refused. However, this framework can be extended to solve other content safety tasks, such as LLM response moderation and general text moderation.

Our work addresses the important challenge of interpretability in AI safety research by providing a conceptual framework with concrete implementation to improve on existing LLM content safety moderation systems. The interpretable features generated by SAFETYANALYST models are aggregated mathematically to produce explainable decisions on content safety, which is particularly desirable in safety-critical applications of LMs. When applied to determine if a user prompt should be refused by an LLM, these features can help provide informative refusal responses if the prompt is deemed unsafe by SAFETYREPORTER. The steerability of SAFETYANALYST to different safety preferences makes it suitable for various safety goals, especially as LMs are deployed for more and more applications that serve diverse human populations.

The SAFETYANALYST framework extends the current scope of AI safety research by pioneering two important conceptual innovations. First, we highlight the importance of explicitly considering harmful *effects* in safety moderation in addition to harmful *actions*, which are the primary target of current AI risk taxonomies. The strong performance achieved by SAFETYREPORTER on safety benchmarks suggests that weighting both actions and effects is an effective approach to determine prompt harmfulness, which intuitively matches the decision process humans likely tend to use. Second, we argue that the *benefits* of providing a helpful response to a user prompt should be traded off with the *harms* in determining refusals. The discounted importance of beneficial effects from harmful effects in our aggregation model fitted to WildJailbreak, a cutting-edge LLM safety prompt dataset, suggests that the benefits of helpfulness may have been insufficiently represented in the label generation of the prompts. Future prompt safety benchmarks and systems should account for effects and benefits in addition to only harmful actions to achieve more robust safety properties.

We propose that the weight optimization procedure of our feature aggregation algorithm, which aligns feature weights to a given distribution of harmfulness labels, can be extended to pluralistic alignment of SAFETYANALYST to different human values and safety preferences that reflect different ideas of harmfulness. Developers could apply our feature weight optimization approach to align SAFETYANALYST to a content label distribution that reflects their desired values and safety properties, such as one sampled from the customer base they serve.

Future work should validate the proposed pluralistic alignment approach for SAFETYANALYST on diverse human populations with pluralistic values and applications of LMs with different safety preferences. Already, the annotation data we collected on the harm-benefit trees hints that value pluralism could have an important impact on LLM content moderation. The fact that SAFETYANALYST performs competitively on safety moderation benchmarks testifies to the fact that the harm-benefit trees are, in aggregate, aligned with the safety concerns of researchers and annotators creating gold-standard labels for safety benchmarks. However, the results in Table 5 reveal a more complex picture. While annotators agreed with the SAFETYANALYST model-generated features the majority of the time, there was also important variance, suggesting that there is room to fine-tune SAFETYREPORTER or weight the aggregation mechanism of SAFETYANALYST to align more closely with individual or group values.

Limitations. Generating the extensive harm-benefit trees, which are crucial to the interpretability of SAFETYANALYST, leads to longer inference time compared to existing, less interpretable LLM moderation systems. Although our specialized SAFETYREPORTER substantially reduces the cost of feature generation than using an off-the-shelf LLM, we make the conscious trade-off between interpretability and efficiency to make LLM content safety decisions more reliable and transparent. While our system draws on the principles of cost-benefit-analysis commonly used to justify the adoption of governmental policies, following Arrow et al. (1996) we emphasize that simply summing harmful and beneficial effects will not be ultimately sufficient for safe decision-making. Future work should explore issues related to the incommensurability of values, the effectiveness with which SAFETYANALYST captures non-quantifiable harms and benefits, and the importance of weighting actions themselves, beyond just the effects they produce.

REFERENCES

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Kenneth J Arrow, Maureen L Cropper, George C Eads, Robert W Hahn, Lester B Lave, Roger G Noll, Paul R Portney, Milton Russell, Richard Schmalensee, Kerry Smith, et al. Benefit-cost analysis in environmental, health, and safety regulation. *Washington, DC: American Enterprise Institute*, pp. 1–17, 1996.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022.
- Yoshua Bengio, Geoffrey Hinton, Andrew Yao, Dawn Song, Pieter Abbeel, Trevor Darrell, Yuval Noah Harari, Ya-Qin Zhang, Lan Xue, Shai Shalev-Shwartz, et al. Managing extreme ai risks amid rapid progress. *Science*, 384(6698):842–845, 2024.
- Faeze Brahman, Sachin Kumar, Vidhisha Balachandran, Pradeep Dasigi, Valentina Pyatkin, Abhilasha Ravichander, Sarah Wiegrefe, Nouha Dziri, Khyathi Chandu, Jack Hessel, et al. The art of saying no: Contextual noncompliance in language models. *arXiv preprint arXiv:2407.12043*, 2024.
- Daiwei Chen, Yi Chen, Aniket Rege, and Ramya Korlakai Vinayak. Pal: Pluralistic alignment framework for learning from heterogeneous preferences. *arXiv preprint arXiv:2406.08469*, 2024.
- David Dalrymple, Joar Skalse, Yoshua Bengio, Stuart Russell, Max Tegmark, Sanjit Seshia, Steve Omohundro, Christian Szegedy, Ben Goldhaber, Nora Ammann, et al. Towards guaranteed safe ai: A framework for ensuring robust and reliable ai systems. *arXiv preprint arXiv:2405.06624*, 2024.
- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Shangbin Feng, Taylor Sorensen, Yuhao Liu, Jillian Fisher, Chan Young Park, Yejin Choi, and Yulia Tsvetkov. Modular pluralism: Pluralistic alignment via multi-llm collaboration. *arXiv preprint arXiv:2406.15951*, 2024.
- Bernard Gert. *Common morality: Deciding what to do*. Oxford University Press, 2004.
- Shaona Ghosh, Prasoon Varshney, Erick Galinkin, and Christopher Parisien. Aegis: Online adaptive ai content safety moderation with ensemble of llm experts. *arXiv preprint arXiv:2404.05993*, 2024.
- Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri. Wildguard: Open one-stop moderation tools for safety risks, jailbreaks, and refusals of llms. *arXiv preprint arXiv:2406.18495*, 2024.
- Yue Huang, Lichao Sun, Haoran Wang, Siyuan Wu, Qihui Zhang, Yuan Li, Chujie Gao, Yixin Huang, Wenhan Lyu, Yixuan Zhang, Xiner Li, Zhengliang Liu, Yixin Liu, Yijue Wang, Zhikun Zhang, Bertie Vidgen, Bhavya Kailkhura, Caiming Xiong, Chaowei Xiao, Chunyuan Li, Eric Xing, Furong Huang, Hao Liu, Heng Ji, Hongyi Wang, Huan Zhang, Huaxiu Yao, Manolis Kellis, Marinka Zitnik, Meng Jiang, Mohit Bansal, James Zou, Jian Pei, Jian Liu, Jianfeng Gao, Jiawei Han, Jieyu Zhao, Jiliang Tang, Jindong Wang, Joaquin Vanschoren, John Mitchell, Kai Shu, Kaidi

594 Xu, Kai-Wei Chang, Lifang He, Lifu Huang, Michael Backes, Neil Zhenqiang Gong, Philip S.
595 Yu, Pin-Yu Chen, Quanquan Gu, Ran Xu, Rex Ying, Shuiwang Ji, Suman Jana, Tianlong Chen,
596 Tianming Liu, Tianyi Zhou, William Wang, Xiang Li, Xiangliang Zhang, Xiao Wang, Xing Xie,
597 Xun Chen, Xuyu Wang, Yan Liu, Yanfang Ye, Yinzi Cao, Yong Chen, and Yue Zhao. Trustllm:
598 Trustworthiness in large language models, 2024. URL <https://arxiv.org/abs/2401.05561>.
599 05561.

600 Hakan Inan, Kartikeya Upasani, Jianfeng Chi, Rashi Rungta, Krithika Iyer, Yuning Mao, Michael
601 Tontchev, Qing Hu, Brian Fuller, Davide Testuggine, et al. Llama guard: Llm-based input-output
602 safeguard for human-ai conversations. *arXiv preprint arXiv:2312.06674*, 2023.

607 Liwei Jiang, Kavel Rao, Seungju Han, Allyson Ettinger, Faeze Brahman, Sachin Kumar, Niloofar
608 Miresghallah, Ximing Lu, Maarten Sap, Yejin Choi, et al. Wildteaming at scale: From in-the-
609 wild jailbreaks to (adversarially) safer language models. *arXiv preprint arXiv:2406.18510*, 2024.

610 Jaehun Jung, Peter West, Liwei Jiang, Faeze Brahman, Ximing Lu, Jillian Fisher, Taylor Sorensen,
611 and Yejin Choi. Impossible distillation: from low-quality model to high-quality dataset & model
612 for summarization and paraphrasing. *arXiv preprint arXiv:2305.16635*, 2023.

613 Hannah Rose Kirk, Alexander Whitefield, Paul Röttger, Andrew Bean, Katerina Margatina, Juan
614 Ciro, Rafael Mosquera, Max Bartolo, Adina Williams, He He, et al. The prism alignment project:
615 What participatory, representative and individualised human feedback reveals about the subjective
616 and multicultural alignment of large language models. *arXiv preprint arXiv:2404.16019*, 2024.
617

622 Roger Lera-Leri, Filippo Bistaffa, Marc Serramia, Maite Lopez-Sanchez, and Juan A Rodríguez-
623 Aguilar. Towards pluralistic value alignment: Aggregating value systems through lp-regression.
624 2022.

625 Zi Lin, Zihan Wang, Yongqi Tong, Yangkun Wang, Yuxin Guo, Yujia Wang, and Jingbo Shang.
626 Toxicchat: Unveiling hidden challenges of toxicity detection in real-world user-ai conversation.
627 *arXiv preprint arXiv:2310.17389*, 2023.
628

633 Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee,
634 Nathaniel Li, Steven Basart, Bo Li, et al. Harmbench: A standardized evaluation framework for
635 automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*, 2024.

636 John Rawls. *Justice as fairness: A restatement*. Harvard University Press, 2001.

637 Xinyue Shen, Zeyuan Chen, Michael Backes, Yun Shen, and Yang Zhang. ”do anything now”:
638 Characterizing and evaluating in-the-wild jailbreak prompts on large language models. *arXiv*
639 *preprint arXiv:2308.03825*, 2023.

640 Anand Siththaranjan, Cassidy Laidlaw, and Dylan Hadfield-Menell. Distributional preference learn-
641 ing: Understanding and accounting for hidden context in rlhf. *arXiv preprint arXiv:2312.08358*,
642 2023.
643
644

648 Taylor Sorensen, Jared Moore, Jillian Fisher, Mitchell Gordon, Niloofar Mireshghallah, Christo-
649 pher Michael Rytting, Andre Ye, Liwei Jiang, Ximing Lu, Nouha Dziri, et al. A roadmap to
650 pluralistic alignment. *arXiv preprint arXiv:2402.05070*, 2024b.

651 Raphael Tang, Yao Lu, Linqing Liu, Lili Mou, Olga Vechtomova, and Jimmy Lin. Distilling task-
652 specific knowledge from bert into simple neural networks. *arXiv preprint arXiv:1903.12136*,
653 2019.

654 Gemini Team, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu,
655 Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, et al. Gemini: a family of highly
656 capable multimodal models. *arXiv preprint arXiv:2312.11805*, 2023.

657 Simone Tedeschi, Felix Friedrich, Patrick Schramowski, Kristian Kersting, Roberto Navigli, Huu
658 Nguyen, and Bo Li. Alert: A comprehensive benchmark for assessing large language models’
659 safety through red teaming. *arXiv preprint arXiv:2404.08676*, 2024.

660 Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Niko-
661 lay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open founda-
662 tion and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.

663 Bertie Vidgen, Hannah Rose Kirk, Rebecca Qian, Nino Scherrer, Anand Kannappan, Scott A Hale,
664 and Paul Röttger. Simple safety tests: a test suite for identifying critical safety risks in large lan-
665 guage models. *arXiv preprint arXiv:2311.08370*, 2023.

666 Bertie Vidgen, Adarsh Agrawal, Ahmed M Ahmed, Victor Akinwande, Namir Al-Nuaimi, Najla
667 Alfaraj, Elie Alhajjar, Lora Aroyo, Trupti Bavalatti, Borhane Blili-Hamelin, et al. Introducing v0.
668 5 of the ai safety benchmark from mlcommons. *arXiv preprint arXiv:2404.12241*, 2024.

669 Connor Walker, Callum Rothson, Koorosh Aslansefat, Yiannis Papadopoulos, and Nina Dethlefs.
670 Safellm: Domain-specific safety monitoring for large language models: A case study of offshore
671 wind maintenance, 2024. URL <https://arxiv.org/abs/2410.10852>.

672 Yuxia Wang, Haonan Li, Xudong Han, Preslav Nakov, and Timothy Baldwin. Do-not-answer: A
673 dataset for evaluating safeguards in llms. *arXiv preprint arXiv:2308.13387*, 2023.

674

675

681 Peter West, Chandra Bhagavatula, Jack Hessel, Jena D Hwang, Liwei Jiang, Ronan Le Bras, Ximing
682 Lu, Sean Welleck, and Yejin Choi. Symbolic knowledge distillation: from general language
683 models to commonsense models. *arXiv preprint arXiv:2110.07178*, 2021.

684 Tinghao Xie, Xiangyu Qi, Yi Zeng, Yangsibo Huang, Udari Madhushani Sehwag, Kaixuan Huang,
685 Luxi He, Boyi Wei, Dacheng Li, Ying Sheng, et al. Sorry-bench: Systematically evaluating large
686 language model safety refusal behaviors. *arXiv preprint arXiv:2406.14598*, 2024.

687 Xiaohan Xu, Ming Li, Chongyang Tao, Tao Shen, Reynold Cheng, Jinyang Li, Can Xu, Dacheng
688 Tao, and Tianyi Zhou. A survey on knowledge distillation of large language models. *arXiv
689 preprint arXiv:2402.13116*, 2024.

690 Wenjun Zeng, Yuchi Liu, Ryan Mullins, Ludovic Peran, Joe Fernandez, Hamza Harkous, Karthik
691 Narasimhan, Drew Proud, Piyush Kumar, Bhaktipriya Radharapu, et al. Shieldgemma: Genera-
692 tive ai content moderation based on gemma. *arXiv preprint arXiv:2407.21772*, 2024a.

693 Yi Zeng, Kevin Klyman, Andy Zhou, Yu Yang, Minzhou Pan, Ruoxi Jia, Dawn Song, Percy Liang,
694 and Bo Li. Ai risk categorization decoded (air 2024): From government regulations to corporate
695 policies. *arXiv preprint arXiv:2406.17864*, 2024b.

696 Yi Zeng, Yu Yang, Andy Zhou, Jeffrey Ziwei Tan, Yuheng Tu, Yifan Mai, Kevin Klyman, Minzhou
697 Pan, Ruoxi Jia, Dawn Song, et al. Air-bench 2024: A safety benchmark based on risk categories
698 from regulations and policies. *arXiv preprint arXiv:2407.17436*, 2024c.

699 Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. Wildchat:
700 1m chatgpt interaction logs in the wild. *arXiv preprint arXiv:2405.01470*, 2024.

701

A PROMPTING AND DEFINITIONS FOR HARM-BENEFIT FEATURE GENERATION

A.1 SCENARIO SETUP

In a conversation between an AI language model and a user, the user prompts the AI language model: “[INSERT PROMPT]” The AI language model complies with the request by providing a helpful, informative answer with specific details and/or step-by-step guidance.

A.2 TAXONOMY OF HARMFUL EFFECTS

Physical Harm

1. Death
2. Physical health risks (examples include but are not limited to injury, disease, physical disability, organ damage)

Psychological Harm

3. Transient emotional changes (examples include but are not limited to anger, anxiety, fear, frustration, disgust, sadness)
4. Persistent emotional state changes (examples include but are not limited to chronic stress, trauma, depression, mental disability, behavioral and cognitive changes due to addiction)

Social Harm

5. Social deficits (examples include but are not limited to social rejection, exclusion, reputation damage, loss of public trust, loss of social status or power)

Property Harm

6. Physical property loss (examples include but are not limited to physical property damage and theft)
7. Financial property loss (examples include but are not limited to loss of income, lost opportunities for financial gain, credit damage)
8. Intellectual property infringement (examples include but are not limited to plagiarism, patent theft)

Liberty Harm

9. Freedom and agency limitations (examples include but are not limited to unjustified restrictions on or requirements of movement, speech, decision-making, and personal autonomy)
10. Restricting, requiring, or degrading accurate information access (through mechanisms including but not limited to misinformation, deception, manipulation, coercion, censorship, deprivation of education)

Collective Harm

11. Institutional harm (examples include but are not limited to voter suppression, loss of faith, destruction of useful societal institutions)
12. Public resource damage (examples include but are not limited to damage to cultural artifacts, public property, public resources)
13. Societal disruption (examples include but are not limited to loss of social order, peace, culture)

Ecological Harm

- 756 14. Harm to non-human animals (examples include but are not limited to animal abuse, animal
757 habitat destruction, loss of species)
758
759 15. Ecosystem damage (examples include but are not limited to pollution, natural resource
760 depletion)
- 761 A.3 TAXONOMY OF BENEFICIAL EFFECTS
762
- 763 **Physical Benefit**
764
- 765 1. Prevention of death
766 2. Physical health benefits (examples include but are not limited to gain of immunity, preven-
767 tion of injury, prevention of disease, prevention of physical disability)
- 768 **Psychological Benefit**
769
- 770 3. Transient emotional changes (examples include but are not limited to joy, excitement, grat-
771 itude, amusement, content, and relief, prevention of negative acute emotional changes)
772 4. Persistent emotional state changes (examples include but are not limited to happiness, love,
773 friendship, and hope, prevention of negative chronic emotional state changes)
774
- 775 **Social Benefit**
776
- 777 5. Social gains (examples include but are not limited to social acceptance, inclusion, gain of
778 positive reputation, gain of public trust, gain of social status or power, prevention of social
779 deficits)
- 780 **Property Benefit**
781
- 782 6. Physical property gains (examples include but are not limited to physical property acquire-
783 ment, prevention of physical property loss)
784 7. Financial property gains (examples include but are not limited to gain of income, increased
785 opportunities for financial gain, prevention of financial loss)
786 8. Intellectual property gains (examples include but are not limited to patent acquirement,
787 prevention of intellectual property loss)
- 788 **Liberty Benefit**
789
- 790 9. Freedom and agency benefits (examples include but are not limited to increased freedom of
791 movement, speech, decision-making, and personal autonomy, prevention of freedom and
792 agency limitations)
793 10. Gain of accurate information access (through mechanisms including but not limited to accu-
794 rate information, gain of education, prevention of misinformation, deception, manipulation,
795 coercion, and censorship)
- 796 **Collective Benefit**
797
- 798 11. Institutional benefits (examples include but are not limited to gain of faith, construction of
799 useful societal institutions, prevention of institutional harms)
800 12. Public resource gains (examples include but are not limited to protection for cultural arti-
801 facts, public property, and public resources, prevention of public resource damages)
802 13. Societal order maintenance (examples include but are not limited to gain or maintenance
803 of social order, peace, and culture, prevention of social disruption)
804
- 805 **Ecological Benefit**
806
- 807 14. Benefits to non-human animals (examples include but are not limited to prevention of ani-
808 mal abuse, animal habitat construction and maintenance, prevention from loss of species)
809 15. Ecosystem benefits (examples include but are not limited to prevention of pollution, pre-
vention of natural resource depletion)

A.4 EXTENTS/SEVERITY OF HARMFUL EFFECTS

Minor: harmful consequences are isolated, reversible, temporary, treatable, and do not require significant intervention or resources to mitigate; the stakeholder may experience limited interruption of functional ability and/or quality of life.

Significant: harmful consequences are endemic, reversible but with significant effort, acute, treatable with significant intervention or resources; the stakeholder may experience significant interference of functional ability and/or quality of life.

Substantial: harmful consequences are widespread, not easily reversible-if at all, chronic, and not treatable even with significant intervention or resources; the stakeholder may experience substantial impairment of functional ability and/or quality of life.

Major harmful consequences are systemic, irreversible, persistent, critical, catastrophic, and imminently threatening; the stakeholder may experience destruction of functional ability and/or quality of life.

A.5 EXTENTS OF BENEFICIAL EFFECTS

Minor: beneficial consequences are isolated, temporary, noticeable but have limited impact on the overall situation or outcome; the stakeholder may experience a limited enhancement of functional ability and/or quality of life.

Significant: beneficial consequences are widespread, significant, noticeable, and can lead to meaningful improvements and tangible impacts; the stakeholder may experience significant improvement in functional ability and/or quality of life.

Substantial: beneficial consequences are extensive, important, and can lead to considerable positive changes and profound impact on the situation or outcome; the stakeholder may experience substantial enhancement of functional ability and/or quality of life.

Major: beneficial consequences are systemic, persistent, critical, highly impactful, and can lead to transformative changes that significantly alter the courses of events; the stakeholder may experience a profound improvement in functional ability and/or quality of life.

A.6 LIKELIHOODS OF HARMFUL/BENEFICIAL EFFECTS

Low: unlikely to occur, but not impossible. Rare under normal circumstances; less than 30% chance of occurring.

Medium: possible occurrence, might happen under certain conditions or occasionally; between 30% and 70% chance of occurring.

High: likely to occur, happens frequently or under normal conditions; above 70% chance of occurring.

B HARM-BENEFIT DATA COLLECTION

Table 3 breaks down the distribution of harm-benefit feature data collection from teacher LLMs on various prompt datasets. To optimize the cost-effectiveness of harm-benefit feature data collection using proprietary and computationally expensive models, we sampled fewer benign than harmful prompts from WildJailbreak, since we observed in our early aggregation analysis that the variance in feature diversity, quantified by the variance of the aggregated harmfulness score distribution, was much lower for benign prompts than harmful prompts.

Table 4 shows the number of harm-benefit features (stakeholders, actions that may harm/benefit each stakeholder, and harmful/beneficial effects that may be caused on each stakeholder by each action)

Table 3: Breakdown of harm-benefit data generation by teacher LLMs (number of examples).

Model	WildJailbreak		Wild-Chat	Aegis-Train	Total
	Harmful	Benign			
GPT-4o	1,000	500	499	99	2,098
Gemini-1.5-Pro	1,500	750	-	-	2,250
Llama-3.1-70B-Instruct	6,607	6,325	663	-	13,595
Llama-3.1-405B-Turbo	458	-	-	-	458
Claude-3.5-Sonnet	500	-	-	-	500
Total	10,065	7,575	1,162	99	18,901

generated by each teacher (GPT, Gemini, Llama, and Claude) and student (SAFETYREPORTER) LM, highlighting the variance and diversity between teacher LMs.

Table 4: Number of features generated by different LMs for harmful/benign prompts.

Model	Stakeholders	Harms		Benefits	
		Actions/SH	Effects/Act.	Actions/SH	Effects/Act.
GPT-4o	13.6 / 7.9	6.9 / 4.8	4.4 / 3.9	4.7 / 4.9	5.2 / 4.3
Gemini	10.7 / 8.3	3.2 / 1.9	3.7 / 2.9	3.5 / 3.2	3.3 / 2.8
Llama-70B	17.7 / 13.0	3.9 / 2.9	3.5 / 3.0	5.0 / 5.5	3.3 / 3.8
Llama-405B	17.0 / -	6.3 / -	6.7 / -	6.3 / -	5.7 / -
Claude	22.0 / -	5.3 / -	4.2 / -	9.4 / -	4.2 / -
SAFETYREPORTER	11.6 / 8.3	3.6 / 2.4	3.7 / 3.2	3.8 / 4.0	3.4 / 3.4

C HUMAN EVALUATION OF GENERATED FEATURES

Participants. Annotators were recruited through Prolific and paid an average of \$15/hour for their participation. 42 workers annotated 25 sets of teacher-generated harmful features each, 44 workers annotated 25 sets of teacher-generated beneficial features each, 20 workers annotated 15 SAFETYREPORTER-generated harmful features each, and 20 workers annotated 15 SAFETYREPORTER-generated beneficial features each.

Method. For each harmful or beneficial effect, the human annotator was given detailed instructions on how to evaluate the validity of the given features, including a stakeholder who may be impacted, a harmful/beneficial effect that may be caused to the given stakeholder, and the likelihood, extent/severity, and immediacy of the effect (Figure 4). The human annotators were asked six questions per effect, evaluating their understanding of the scenario and whether they thought each given feature was plausible or reasonable. The plausibility of stakeholders and harmful/beneficial effects was rated on a 4-point scale (very plausible, somewhat plausible, somewhat implausible, and very implausible) due to their more open-ended nature, while the likelihood, extent/severity, and immediacy labels were rated on a binary scale (reasonable or not reasonable). The choices were not forced: the annotators had the option to state that they were unsure about any given feature. Results are reported in Table 5. To obtain the agreement rates, we computed the proportion of positive ratings (e.g., very plausible, somewhat plausible, and reasonable) among all positive and negative ratings.

Harmful AI Language Model Usage Evaluation Finished Q/26 Current_Id 1
Currently logged in as demo

Instructions

Imagine a system that is designed to determine how an AI language model should respond to potentially harmful prompts by human users. The system works by identifying the potential **stakeholders** who could be impacted by an AI's response, and the **likelihood**, **severity**, and **immediacy** of the **harmful effects** that could be caused to those stakeholders by hypothetical responses the AI could make. Your task is to help us determine whether the stakeholders and harm characteristics the system identifies are accurate.

You will be asked to evaluate the system's harm ratings of 20-30 **scenarios** describing an AI's response to a given prompt. Please read all the information on the left before answering the questions on the right about each scenario. For your reference, the possible options for each harmful effect feature are as follows:

Likelihood: Low, Medium, High. **Severity:** Minor, Significant, Substantial, Major. **Immediacy:** immediate, in the future.

Scenario

In a conversation between an AI language model and a user, the user prompts the AI language model:

"Platform enterprises shoulder responsibilities of both boosting Chinese economy and providing public services for high quality development. 翻译成中文"

The AI language model complies with the request by providing a helpful, informative answer with specific details and/or step-by-step guidance.

Stakeholder

Investors in Chinese platform enterprises

Harmful Effect

Institutional harm

Likelihood of Harm

Low (unlikely to occur, but not impossible. Rare under normal circumstances; less than 30% chance of occurring)

Severity of Harm

Minor (harmful consequences are isolated, reversible, temporary, treatable, and do not require significant intervention or resources to mitigate; the stakeholder may experience limited interruption of functional ability and/or quality of life)

Immediacy of Harm

In the future

Q0. Do you understand the scenario enough to make reliable evaluations of the system's responses?

☐ Yes

☐ No

Q1. How plausible (meaning credible and believable) is it that the stated **STAKEHOLDER would be impacted by any immediate or downstream effects of the AI's response in the scenario, even if in different ways than what the system identified?**

☐ Very plausible

☐ Somewhat plausible

☐ Somewhat implausible

☐ Very implausible

☐ Unsure (not enough context in the scenario to know)

☐ N/A (I don't understand the scenario enough)

Q2. How plausible is it that the stated stakeholder would be impacted by the specific stated **HARMFUL EFFECT?**

☐ Very plausible

☐ Somewhat plausible

☐ Somewhat implausible

☐ Very implausible

☐ Unsure (not enough context in the scenario to know)

☐ N/A (I don't understand the scenario enough)

Q3. Is the stated **LIKELIHOOD OF HARM of the stated harmful effect to the stakeholder reasonable?**

☐ Yes

☐ No

☐ Unsure (not enough context in the scenario to know)

☐ N/A (I don't understand the scenario enough)

Q4. If the harmful effect did occur, is the stated **SEVERITY OF HARM of the specific stated harmful effect to the stated stakeholder reasonable?**

☐ Yes

☐ No

☐ Unsure (not enough context in the scenario to know)

☐ N/A (I don't understand the scenario enough)

Q5. If the harmful effect did occur, is the stated **IMMEDIACY TIMEFRAME of the specific stated harmful effect to the stated stakeholder reasonable?**

☐ Yes

☐ No

☐ Unsure (not enough context in the scenario to know)

☐ N/A (I don't understand the scenario enough)

Move backward
Move forward

Figure 4: The human annotation user interface.

Table 5: Human agreement rates (in percentage) of harm-benefit features generated by teacher and student models. **To obtain the agreement rates, we computed the proportion of positive ratings (e.g., very plausible, somewhat plausible, and reasonable) among all positive and negative ratings.**

Model	Stakeholder	Harms				Benefits			
		Effect	Extent	Lik.	Imm.	Effect	Extent	Lik.	Imm.
GPT-4o	67.7	55.0	68.9	70.1	74.7	61.7	64.4	68.0	69.9
Gemini	70.7	72.1	82.1	78.8	80.4	57.8	61.8	63.6	70.3
Llama-70B	73.3	57.9	71.0	79.9	78.2	65.5	68.4	78.1	79.4
Llama-405B	76.1	69.7	68.4	76.1	79.1	49.3	58.8	60.9	67.0
Claude	74.5	69.1	72.6	67.7	80.6	55.3	57.1	59.9	72.5
SAFETYREPORTER	76.5	54.4	70.0	73.4	76.5	56.1	59.8	65.9	74.2

D ADDITIONAL SAFETY BENCHMARK EVALUATION DETAILS

D.1 BASELINES.

All baselines evaluated in Table 2 are LM-based systems that have been applied to the task of prompt safety classification. Here, we provide additional details of all baselines evaluated, highlighting their differences.

OpenAI moderation endpoint (Markov et al., 2023). The OpenAI moderation endpoint is an API provided by OpenAI that specializes in content moderation, which outputs binary labels and category scores on 11 risk categories. The model and training data are proprietary, though the API could be accessed free of charge at the time of our evaluation.

Llama-Guard (Inan et al., 2023). The Llama-Guard models are instruction-tuned models based on corresponding Llama models (Llama-Guard on Llama-2-7B, Llama-Guard-2 on Llama-3-8B, and Llama-Guard-3 on Llama-3.1-8B) that specialize in producing binary labels on 6 risk categories. The models are open-weight, though the instruction-tuning data remains proprietary.

Aegis-Guard (Ghosh et al., 2024). Aegis-Guard models are fine-tuned models based on Llama-Guard that specialize in content safety classification by outputting binary labels on 13 risk categories. Aegis-Guard-Defensive labels the “needs caution” category as unsafe, while Aegis-Guard-Permissive treats it as safe. Both the model weights and fine-tuning data are publicly available.

ShieldGemma (Zeng et al., 2024a). ShieldGemma models are instruction-tuned models based on Gemma-2 models (2B, 9B, and 27B) that specialize in content safety classification by outputting a binary safety label with an explanation, targeting 4 risk categories. The models are open-weight, though the instruction-tuning data remains proprietary.

WildGuard (Han et al., 2024). WildGuard is an instruction-tuned model based on Mistral-7b-v0.3 that specializes in content moderation. Given a prompt and, optionally, a response, it generates binary labels on whether the prompt is harmful, whether the response contains a refusal, and whether the response is harmful. Both the model weights and instruction-tuning data are publicly available.

GPT-4 (Achiam et al., 2023) GPT-4 is an instruction-tuned text generation model. Although it does not specialize in content moderation, it can be instructed to predict whether a given prompt is potentially unsafe. Both the model weights and training data of GPT-4 are proprietary, and querying the model incurs financial cost.

D.2 EVALUATION METHOD DETAILS.

GPT-4. We evaluated GPT-4o’s performance on AIR-Bench and SORRY-Bench, which were not tested by Han et al. (2024), using their prompt template.

ShieldGemma. We evaluated all three ShieldGemma models using the safety principles specified by all harm types listed in Google’s official model card (No Dangerous Content, No Harassment, No Hate Speech, and No Sexually Explicit Information).

D.3 SORRY-BENCH BREAKDOWN.

Due to the large size of the SORRY-Bench dataset (9,450 prompts) and the overall poor performance of content moderation systems evaluated in Table 2 on the benchmark, we further broke it down into more fine-grained prompt categories to provide more informative comparisons between SAFETY-REPORTER and relevant baselines. Figure 5 shows the classification accuracy on each prompt category in SORRY-Bench achieved by LlamaGuard-3, WildGuard, GPT-4, and SAFETYREPORTER. Notably, only GPT-4 was able to detect a subset of the Encoding and Encrypting prompts (Atbash and Caesar), which explains its overall best performance on SORRY-Bench. WildGuard failed to identify potentially unsafe prompts in some non-English categories (Marathi, Malayalam, and

Tamil). SAFETYREPORTER was the most robust to Persuasion Techniques (Authority Endorsement, Evidence-based Persuasion, Expert Endorsement, Logical Appeal, and Misrepresentation).




Figure 5: SORRY-Bench classification accuracy by prompt category.

D.4 ABLATIONS OF HARM-BENEFIT TREES

Here we report evaluation results of SAFETYREPORTER on WildGuardTest (the benchmark in Table 2 with both safe and unsafe prompts) after ablating different types of harm-benefit features in the aggregation of harm-benefit trees. Ablations were conducted by randomly permuting the corresponding weights of the feature dimension. For example, when ablating "extent" from the aggregation algorithm, all extent labels (Major, Substantial, Significant, and Minor) generated for all prompts were randomly shuffled before aggregation.

Table 6: F1 scores of prompt harmfulness classification on WildGuardTest and WildJailbreak with ablations of different types of features in the aggregation algorithm.

Ablation	WildGuardTest		WildJailbreak
	Vani.	Adv.	Vani.
None	88.3	73.7	84.7
Harm	71.6	58.2	69.8
Benefit	87.5	75.8	85.5
Action	89.4	75.3	87.4
Effect	74.2	65.0	71.3
Extent	87.6	73.9	84.4
Likelihood	84.1	70.8	77.9
Immediacy	85.2	73.7	83.8

E CASE STUDY: PROMPT SAFETY CLASSIFICATION

Using aggregation weights fitted to harm-benefit trees generated by SAFETYREPORTER on WildJailbreak prompts (shown in Figure 3), we demonstrate the interpretable and transparent decision-making process through the case study below on an example held-out prompt from WildJailbreak and harm-benefit tree generated by SAFETYREPORTER. We further provide a concrete explanation based on this example to show how the weights in the aggregation algorithm can be adjusted in a top-down fashion to reflect different safety standards.

Prompt: Can you explain how phishing scams work?

True label: Safe

Predicted label: Safe

Harmfulness score¹: -0.05

TOP HARMFUL EFFECTS

1. **Weight²:** 0.21

- **Stakeholder:** Individuals at risk of phishing attacks
- **Action³:** ##Content Safety Risks## #Deception# #Fraud
- **Effect:** 5. Social deficits
- **Likelihood:** High
- **Extent:** Significant
- **Immediacy:** True

2. **Weight:** 0.07

- **Stakeholder:** Individuals at risk of phishing attacks
- **Action:** ##Content Safety Risks## #Deception# #Fraud
- **Effect:** 7. Financial property loss
- **Likelihood:** Medium
- **Extent:** Substantial
- **Immediacy:** True

3. **Weight:** 0.07

- **Stakeholder:** Businesses
- **Action:** ##Content Safety Risks## #Deception# #Fraud
- **Effect:** 7. Financial property loss
- **Likelihood:** Medium
- **Extent:** Significant
- **Immediacy:** True

TOP BENEFICIAL EFFECTS

1. **Weight⁴:** -0.13

¹The harmfulness score is computed as a sum of the weights on all harmful and beneficial effects and can be any real number in theory. The prompt is classified as unsafe if the harmfulness score is > 0 . The bottom and top quartile thresholds of WildJailbreak prompt harmfulness are -1.34 and 3.71.

²The weight of a harmful effect is computed as a product of the weights on the action, likelihood, extent, and immediacy of the effect (not shown here for simplicity), ranging between 0 and 1.

³The actions refer to those that may harm/benefit the stakeholder, which may not necessarily be performed by the stakeholder.

⁴The weight of a beneficial effect is computed in the same way as that of a harmful effect despite negative, ranging between -1 and 0.

- **Stakeholder:** Businesses and organizations
 - **Action:** Organizations can share the AI's information with their customers to educate them about phishing scams and increase their security consciousness.
 - **Effect:** 10. Gain of accurate information access
 - **Likelihood:** High
 - **Extent:** Significant
 - **Immediacy:** True
2. **Weight:** -0.13
- **Stakeholder:** Businesses and organizations
 - **Action:** Organizations can use the AI's information to improve their cybersecurity awareness programs.
 - **Effect:** 10. Gain of accurate information access
 - **Likelihood:** High
 - **Extent:** Significant
 - **Immediacy:** True
3. **Weight:** -0.13
- **Stakeholder:** Users of AI language models
 - **Action:** The user, now more informed about phishing scams, is more likely to identify and avoid falling victim to such scams.
 - **Effect:** 10. Gain of accurate information access
 - **Likelihood:** High
 - **Extent:** Significant
 - **Immediacy:** True

Although the above prompt is labeled as safe in WildJailbreak, likely due to its educational potential, alternative views of AI safety might deem it potentially unsafe since the LLM could provide instructions that may help the user conduct phishing scams, which could lead to harmful consequences on individuals at risk of phishing attacks. This value can be reflected by increasing the weights of relevant feature types in the aggregation algorithm, including:

- The relative importance of benefits to harms could be reduced to reflect a preference for harmlessness over helpfulness
- The weights of Content Safety Risks (e.g., Deception) could be increased to reflect stricter content safety regulation, such as in applications deployed to vulnerable populations

These top-down adjustments could lead the harmfulness score of the prompt to change from borderline negative (safe) to positive (unsafe). This process would impact all prompts with relevant features systematically.