
Enforcing boundary conditions for physics-informed neural operators

Abstract

Machine-learning based techniques like physics-informed neural networks (PINNs) and physics-informed neural operators (PINO) are becoming increasingly adept at solving even complex systems of partial differential equations (PDEs). Boundary conditions can be enforced either weakly by penalizing deviations in the loss function or strongly by training a solution structure that inherently matches the prescribed values and derivatives. The former approach is easy to implement but the latter can provide benefits with respect to accuracy and training times. However, previous approaches to strongly enforcing Neumann or Robin boundary conditions require a domain with a fully C^1 boundary and, as we demonstrate, can lead to instability if those boundary conditions are posed on a segment of the boundary that is piecewise C^1 but only C^0 globally. We introduce a generalization of the approach by Sukumar, et al. (2021) and a new approach based on orthogonal projections that overcome this limitation. The performance of these new techniques is compared against weakly and semi-weakly enforced boundary conditions for the scalar Darcy flow equation and the stationary Navier-Stokes equations.

1 Introduction and related work

Various machine learning-based techniques have been applied successfully to solve systems of partial differential equations (PDEs). Most prominently, physics-informed neural networks Raissi, et al. (2019) and many variants and extensions Raissi, et al. (2024) learn the solution to a PDE. They use (often dense) neural networks combined with loss functions that include the PDE residuals to promote physically meaningful solutions and reduce the amount of training data required. Neural operators Li, et al. (2020a); Lu, et al. (2021a) were developed to learn solution operators of PDEs. Based on the Fourier neural operator (FNO) by Li, et al. (2020b), a physics-informed neural operator (PINO) was proposed by Li, et al. (2021). Similar to PINNs it was designed to approximate the solution operator of a parametric partial differential equation by minimizing a residual given by the differential equation instead of training solely on labeled training data. In the following we utilize the FNO framework even though it requires a rectangular domain with a uniform mesh to use the Fast Fourier Transform (FFT). To work on more complex geometries we follow the approach by Lu, et al. (2021c) and choose the minimum bounding box of the underlying domain as computational domain. Alternatively, for more efficient approaches one could utilize the geo-FNO framework proposed by Li, et al. (2022) or geometry-informed neural operators Li, et al. (2023).

Boundary conditions. In physics-informed machine learning, boundary conditions can be enforced in two ways. One is to weakly enforce them by adding a residual term that punishes but does not prohibit differences to the prescribed values. An extension are penalty methods to treat boundary conditions as hard constraints in the optimization Lu, et al. (2021b). However, Toscano, et al. (2025); Zeinhofer, et al. (2024) show that these approaches weaken the decay of the generalization

error. The alternative is to strongly enforce boundary conditions by constructing the solution in a way that it exactly satisfies the boundary conditions. While this is straightforward for Dirichlet boundary conditions Berrone, et al. (2023); Toscano, et al. (2025), Neumann- or Robin boundary conditions are harder to treat. Techniques to do this include Fourier feature embeddings Straub, et al. (2025) or solution structures based on trial solutions using distance functions Manavi, et al. (2024); McFall, et al. (2009). Based on the latter idea, Sukumar, et al. (2021) introduce a flexible method to strongly enforce boundary conditions and observe that it can improve accuracy for PINNs. We will show that their approach is suitable for a certain class of boundary conditions but fails when Neumann conditions are prescribed on a boundary that is not C^1 . They utilize the theory of R-functions by Rvachev, et al. (1995) and approximate distance functions to train solutions to boundary value problems using PINNs. An approximate distance function satisfies two properties.

Definition 1.1 (Distance function). Let $\Gamma \subset \partial\Omega$ be a boundary section. We call a function $\phi : \bar{\Omega} \rightarrow \mathbb{R}$ a *distance function* to Γ if it is zero on Γ and positive in $\bar{\Omega} \setminus \Gamma$.

The other definition is according to Rvachev, et al. (1995).

Definition 1.2 (Normalized function). Let $\Gamma \subset \partial\Omega$ be a boundary section. We call a function $\bar{\phi} : \bar{\Omega} \rightarrow \mathbb{R}$ a *normalized function with respect to Γ* if it satisfies $\bar{\phi} \equiv 0$ and $\frac{\partial \bar{\phi}}{\partial \nu} \equiv 1$ on Γ .

Here, ν denotes the inward pointing normal vector on $\partial\Omega$.

Definition 1.3 (Approximate distance function). Let $\Gamma \subset \partial\Omega$ be a boundary section. We call a function $\phi : \bar{\Omega} \rightarrow \mathbb{R}$ an *approximate distance function* to Γ , if ϕ is both a distance function to Γ in the sense of Definition 1.1 and normalized with respect to Γ in the sense of Definition 1.2.

As pointed out by Sukumar, et al. (2021), suitable approximate distance functions should be C^1 . Otherwise, the Laplacian of the distance function is unbounded at points where the distance function is only C^0 , which causes issue when solving second-order differential equations. While the exact distance function $d(\mathbf{x}) := \min_{\tilde{\mathbf{x}} \in \partial\Omega} \|\mathbf{x} - \tilde{\mathbf{x}}\|$ is an approximate distance function to $\partial\Omega$, it is in general not C^1 . Sukumar, et al. (2021) discuss how to construct approximate distance functions to boundaries that consist of piecewise linear segments which are only C^0 globally. Their approximate distance functions are C^1 in the interior but not at the joining points of the segments. This is not problem for the Dirichlet conditions or Neumann boundary conditions on the annulus they consider, where the boundary is globally C^1 . However, as pointed out by Gladstone, et al. (2022), the non-differentiability of the approximate distance function becomes an issue for Neumann or Robin boundary value on boundaries that are not globally C^1 . We provide a summary of Sukumar, et al. (2021)’s approach in Appendix A.

To illustrate the issue that can arise if $\partial\Omega \notin C^1$, consider a simple Poisson problem with homogeneous Neumann boundary condition

$$\forall (x, y) \in (0, 1)^2 : \quad -\Delta u(x, y) = 2\pi^2 \cos(\pi x) \cos(\pi y). \quad (1)$$

Figure 1 shows resulting solution (middle) as well as the analytical solution (left) and the solution using our generalized approach presented in this paper (right) for comparison. The issue in the middle figure stems from the emergence of instabilities in the corner of the Laplacian of the approximate distance function, cf. Sukumar, et al. (2021, Figure 27).

Contributions and structure of the paper. We propose two novel approaches to prescribe Robin or Neumann boundary conditions on boundaries that are piecewise C^1 but only C^0 globally. First we describe a generalization of the method by Sukumar, et al. (2021) called *generalized local solution structures* or GLSS. The second approach is based on *orthogonal projections* and we refer to it as OP. While it requires certain assumptions on the shape of the boundary, it has fewer unknown functions that need to be learned by the network. GLSS and OP are described in § 2.1 for scalar PDEs and in § 2.2 for systems of PDEs. § 3.1 compares GLSS, OP as well as weakly and semi-weakly boundary conditions for a scalar PDE, the Darcy flow equation. Finally, § 3.2 compares their performance for a standard benchmark from computational fluid dynamics by Turek, et al. (1996) that requires solving the stationary Navier-Stokes equations to model flow around a cylinder.

Limitations. The two new approaches come with some limitations and drawbacks. First, they increase complexity of implementation compared to weakly enforced boundary conditions, in particular if the number of C^1 -segments that form the boundary is high. Second, while the size of the network




Figure 1: Analytical solution (left) and solution with Sukumar, et al. (2021)’s (middle) and our generalized approach (right) to strongly enforcing boundary conditions.

as well as training times change only marginally, we do see an increase of inference times of up to 30%. Third, the approaches are only tested for $\Omega \subset \mathbb{R}^2$, although generalization to 3D should be possible. And lastly, the OP approach only works for boundaries that can be decomposed into segments where each of these segments lies in a hyperplane.

2 Methodology

Let $\Omega \subset \mathbb{R}^2$ be a computational domain, $\mathcal{U}, \mathcal{V}$ Banach spaces, and let $\mathcal{A} \subset \mathcal{V}$ be a set of parameter functions. For a $\mathbf{a} \in \mathcal{A}$ we want to find the solution $\mathbf{u} \in \mathcal{U}$ to the boundary value problem

$$\forall \mathbf{x} \in \Omega : \quad \mathcal{P}(\mathbf{u}(\mathbf{x}), \mathbf{a}(\mathbf{x})) = 0, \quad (2a)$$

$$\forall \mathbf{x} \in \partial\Omega : \quad \mathcal{B}(\mathbf{u}(\mathbf{x}), \mathbf{a}(\mathbf{x})) = 0, \quad (2b)$$

where $\mathcal{P}$ is a differential operator and $\mathcal{B}$ is a boundary condition operator. The idea by Sukumar, et al. (2021) and Rvachev, et al. (1995) is to train suitable functions Ψ_i such that $\tilde{\mathbf{u}}(\Psi_1, \dots, \Psi_I)$ minimizes the PDE-loss from (2a) and, by construction, satisfies the boundary condition exactly.

Definition 2.1 (Solution structure). We call $\tilde{\mathbf{u}} : \{\bar{\Omega} \rightarrow \mathbb{R}^I\} \rightarrow \{\bar{\Omega} \rightarrow \mathbb{R}^n\}$ a solution structure, if $\tilde{\mathbf{u}}(\Psi_1, \dots, \Psi_I)$ satisfies (2b) for any differentiable functions $\Psi_1, \dots, \Psi_I : \bar{\Omega} \rightarrow \mathbb{R}$.

Physics-informed neural networks (PINNs) and physics-informed neural operators (PINO).

For a PINO, the aim is to learn the solution operator $\mathcal{G}_\theta : \mathcal{A} \rightarrow \mathcal{U}$ with $\mathcal{G}_\theta(\mathbf{a}) = \mathbf{u}$ using a set of training parameters $\mathcal{A}_{\text{train}} \subset \mathcal{A}$. We denote the trainable parameters of the neural operator as θ . Note that Sukumar, et al. (2021) present their approach in the context of physics-informed neural network (PINNs) whereas we consider physics-informed neural operators (PINO). However, in the notation above, a PINN learns $u_\theta(\mathbf{a}^*) = G_\theta(\mathbf{a}^*)$ for a *fixed* parameter $\mathbf{a}^*$. That is, it learns one specific instance $G_\theta(\mathbf{a}^*)$ solving the boundary problem (2a). By contrast, a PINO trains on a large set of parameters to learn to mapping $\mathbf{a} \mapsto G_\theta(\mathbf{a})$. The learned solution operator can be further refined to compute the solution for a specific $\mathbf{a}^*$ through continued training only on this parameter (*finetuning*).

In the following, we use the FNO-PINO architecture to learn not a mapping to the PDE solution directly, but to the unknown functions in the *solution structure* (Def. 2.1). With slight abuse of notation we will refer to the output of the neural networks still as $\mathcal{G}_\theta(\mathbf{a})$. Our approaches to enforce boundary conditions can be used for either PINN or PINO. To illustrate applicability to PINNs, in addition to the regular training and finetuning of the PINO, we consider *PINN-like training*, i. e. learning the solution structure for a specific parameter $\mathbf{a}^*$ without training the solution operator first. This approach is essentially a PINN that uses the FNO architecture instead of dense layers. As we focus on boundary conditions, we do not consider data mismatch terms in the loss, but only physics losses. There are three ways to ensure that the trained solution $\mathbf{u}$ satisfies the boundary condition (2).

Weak boundary conditions. Here, the neural operator outputs the solution directly, that is $\mathbf{u} = \mathcal{G}_\theta(\mathbf{a})$, and satisfying the boundary conditions has to be learned during training. This corresponds to

120 solving the minimization problem

$$\min_{\theta} \sum_{\mathbf{a} \in \mathcal{A}_{\text{train}}} \left(\int_{\Omega} \mathcal{P}((\mathcal{G}_{\theta}(\mathbf{a}))(\mathbf{x}), \mathbf{a}(\mathbf{x}))^2 d\mathbf{x} + \int_{\partial\Omega} \mathcal{B}((\mathcal{G}_{\theta}(\mathbf{a}))(\mathbf{x}), \mathbf{a}(\mathbf{x}))^2 dS(\mathbf{x}) \right). \quad (3)$$

121 **Exact boundary conditions.** Here we construct a solution structure in the sense of Definition 2.1
 122 and train the functions Ψ_i such that the parameters satisfy

$$\min_{\theta} \sum_{\mathbf{a} \in \mathcal{A}_{\text{train}}} \int_{\Omega} \mathcal{P}(\tilde{\mathbf{u}}(\mathcal{G}_{\theta}(\mathbf{a}))(\mathbf{x}), \mathbf{a}(\mathbf{x}))^2 d\mathbf{x}. \quad (4)$$

123 for the resulting $\tilde{\mathbf{u}}(\Psi_1, \dots, \Psi_I)$.

124 **Semi-weak / semi-exact boundary conditions.** Lastly we consider a solution structure that satisfies
 125 some boundary conditions but not all. Let $\tilde{\mathcal{B}}$ represent the boundary conditions that $\tilde{\mathbf{u}}(\Psi_1, \dots, \Psi_I)$
 126 does not automatically satisfy. The minimization problem solved in training is then posed as a
 127 combination of the two previous ones

$$\min_{\theta} \sum_{\mathbf{a} \in \mathcal{A}_{\text{train}}} \left(\int_{\Omega} \mathcal{P}(\tilde{\mathbf{u}}(\mathcal{G}_{\theta}(\mathbf{a}))(\mathbf{x}), \mathbf{a}(\mathbf{x}))^2 d\mathbf{x} + \int_{\partial\Omega} \tilde{\mathcal{B}}((\mathcal{G}_{\theta}(\mathbf{a}))(\mathbf{x}), \mathbf{a}(\mathbf{x}))^2 dS(\mathbf{x}) \right). \quad (5)$$

128 We use this approach to enforce Dirichlet boundary conditions exactly and Robin conditions weakly.

129 2.1 Exact boundary conditions for scalar differential equations

130 Starting from Sukumar, et al. (2021)'s approach, the solution consists of two parts, the transfinite
 131 interpolant by Rvachev, et al. (2001) for the boundary and a remainder term in the domain

$$u(\mathbf{x}) = \underbrace{\sum_{i=1}^M w_i(\mathbf{x}) u_i(\mathbf{x})}_{\text{boundary}} + \underbrace{\Psi(\mathbf{x}) \prod_{i=1}^M \phi_i(\mathbf{x})^{\mu_i}}_{\text{domain}}, \quad (6)$$

$$\forall i \in \{1, \dots, M\} : w_i(\mathbf{x}) = \frac{\prod_{j=1, j \neq i}^M \phi_j(\mathbf{x})^{\mu_j}}{\sum_{k=1}^M \prod_{j=1, j \neq i}^M \phi_j(\mathbf{x})^{\mu_j}}, \quad (7)$$

132 where ϕ_i is the distance function to Γ_i . We set $\mu_i = 1$ if $i \in I_D$ and $\mu_i = 2$ if $i \in I_R$, where
 133 $I_D \cup I_R = \{1, \dots, M\}$ are index sets indicating on which segments Dirichlet or Robin boundary
 134 conditions are prescribed.

135 Sukumar, et al. (2021) require an approximate distance function ϕ_i in (6), (7) that is both a distance
 136 function to Γ_i in the sense of Definition 1.1 and normalized with respect to Γ_i in the sense of
 137 Definition 1.2. However, it is not always possible to find such a function ϕ_i that is C^1 everywhere.
 138 Therefore, we propose to use two different functions instead. The function in (6) and (7), which we
 139 still denote as ϕ_i , is only required to be a distance function to Γ_i . The function in the local solution
 140 structures u_i , which we now denote as $\bar{\phi}_i$, only needs to be a normalized function with respect to
 141 Γ_i . Below, we show two ways to choose the local solution structures u_i . For comparison, (35) in
 142 Appendix A shows Sukumar, et al. (2021)'s choice.

143 2.1.1 Generalized local solution structure (GLSS) for piecewise C^1 boundary

144 For pairwise disjoint boundary segments $\Gamma_1, \dots, \Gamma_M$ we use the local solution structure

$$\forall i \in I_D : u_i(\mathbf{x}) = \begin{cases} g_i(\mathbf{x}) + \bar{\phi}_i(\mathbf{x}) \tilde{\Psi}_i(\mathbf{x}), & \phi_i \text{ has a vanishing gradient,} \\ g_i(\mathbf{x}), & \text{else,} \end{cases} \quad (8)$$

$$\forall i \in I_R : u_i(\mathbf{x}) = \Psi_i(\mathbf{x}) - \bar{\phi}_i(\mathbf{x}) \nabla \bar{\phi}_i(\mathbf{x}) \cdot \nabla \Psi_i(\mathbf{x}) + \bar{\phi}_i(\mathbf{x}) (c_i(\mathbf{x}) \Psi_i(\mathbf{x}) - h_i(\mathbf{x})), \quad (9)$$

145 where the $\bar{\phi}_i$ are normalized with respect to Γ_i and the $\tilde{\Psi}_i$ and Ψ_i are unknown functions to be
 146 learned. The difference between (8) and (35) is the term $\bar{\phi}_i(\mathbf{x}) \tilde{\Psi}_i(\mathbf{x})$. Without it, if ϕ_i has a vanishing
 147 gradient, we would prescribe both $u_i = g_i$ on Γ_i and $\frac{\partial u}{\partial \mathbf{n}} = \frac{\partial g_i}{\partial \mathbf{n}}$ on Γ_i , which would overdetermine




Figure 2: Sketch of intersecting boundary segments.

the problem. The additional term $\bar{\phi}_i(\mathbf{x})\tilde{\Psi}_i(\mathbf{x})$ avoids that by introducing another unknown function to be trained. In Sukumar, et al. (2021)'s approach, the ϕ_i were required to be normalized with respect to Γ_i and thus could not have a vanishing gradient, therefore this problem did not arise.

Intersecting boundary segments as sketched in Figure 2 require further modifications to (9), see Appendix B for the reasons. Let Γ_i be a boundary segment with neighbors Γ_a and Γ_b , where A and B are the intersection points. If Γ_i is a Robin segment, we define the function Ψ_i as

$$\Psi_i(\mathbf{x}) = \frac{\phi_B(\mathbf{x})}{\phi_A(\mathbf{x}) + \phi_B(\mathbf{x})}u_A(\mathbf{x}) + \frac{\phi_A(\mathbf{x})}{\phi_A(\mathbf{x}) + \phi_B(\mathbf{x})}u_B(\mathbf{x}) + \phi_A(\mathbf{x})\phi_B(\mathbf{x})\bar{\Psi}_i(\mathbf{x}), \quad (10)$$

where ϕ_A and ϕ_B are distance functions to A and B . The functions u_A and u_B have to be chosen according to type of boundary conditions prescribed on Γ_a and Γ_b . If Γ_a or Γ_b is a Dirichlet segment, we set $u_A = g_a$ or $u_B = g_b$. In case of a Robin segment, we introduce a new unknown function Ψ_A or Ψ_B and define $u_A = \Psi_A$ or $u_B = \Psi_B$. The term $\phi_A(\mathbf{x})\phi_B(\mathbf{x})\bar{\Psi}_i(\mathbf{x})$ only needs to be included if both Γ_a and Γ_b are Dirichlet segments. Function $\bar{\Psi}_i$ is another unknown to be approximated. We show the algorithm for determining local solution structures in the Appendix C].

2.1.2 Orthogonal projections (OP)

If all boundary sections with Robin conditions lie in hyperplanes, i.e., for every $\mathbf{x} \in \Gamma_i$ the normal vectors $\mathbf{n}(\mathbf{x})$ are identical, we can use a simpler approach. We choose the normalized functions $\bar{\phi}_i$ to be the exact signed distance function to the hyperplane containing Γ_i . The local solution structures for Robin conditions are set to

$$\forall i \in I_R : \quad u_i(\mathbf{x}) = \underbrace{\Psi_i(\mathcal{N}(\mathbf{x}; \bar{\phi}_i))}_{\text{boundary value}} + \underbrace{\bar{\phi}_i(\mathbf{x})f_i(\mathcal{N}(\mathbf{x}; \bar{\phi}_i))}_{\text{boundary derivative}}, \quad (11)$$

$$\forall i \in I_R : \quad f_i(\mathbf{x}) = c_i(\mathbf{x})\Psi_i(\mathbf{x}) - h_i(\mathbf{x}) \quad \text{with} \quad \mathcal{N}(\mathbf{x}; \bar{\phi}) := \mathbf{x} - \bar{\phi}(\mathbf{x})\nabla\bar{\phi}(\mathbf{x}) \quad (12)$$

The ansatz uses as a generalized Taylor polynomial expansion by Rvachev, et al. (1995). Here, Ψ_i represents the value and f_i the normal derivative of u_i on Γ_i . The concatenations $\Psi_i \circ \mathcal{N}(\cdot; \bar{\phi}_i)$ and $f_i \circ \mathcal{N}(\cdot; \bar{\phi}_i)$ are the so called normalizers of Ψ_i and f_i , see Appendix D.

The first term in (11) determines the value on Γ_i but has zero derivative whilst the second term is zero on Γ_i but has non-zero derivative equal to the Robin condition. Because $\bar{\phi}_i$ is the exact signed distance function, $\mathcal{N}(\cdot; \bar{\phi}_i)$ is an orthogonal projection mapping its argument onto the corresponding hyperplane. Therefore, each Ψ_i is only evaluated on its corresponding hyperplane and we can set

$$\forall i \in I_R : \quad \Psi_i(\mathbf{x}) = g(\mathbf{x}) + \tilde{\Psi}(\mathbf{x}) \prod_{k \in I_D} \phi_k(\mathbf{x}), \quad (13)$$

where g is a function satisfying all Dirichlet conditions. This avoids discontinuities of the transfinite interpolant at intersection points since only the single function $\tilde{\Psi}$ needs to be trained.

2.2 Exact boundary conditions for systems of partial differential equations

Consider a system of differential equations with solution $\mathbf{u} : \mathbb{R}^2 \supset \Omega \rightarrow \mathbb{R}^n$ and boundary conditions prescribed on segments $\Gamma_1, \dots, \Gamma_M \subset \partial\Omega$ with each Γ_i being C^1 . For $i \in \{1, \dots, M\}$ and $\mathbf{x} \in \Gamma_i$ let $\mathbf{b}_i^{(1)}(\mathbf{x}), \dots, \mathbf{b}_i^{(n)}(\mathbf{x})$ be a basis of $\mathbb{R}^n$. Let $IJ_D, IJ_R \subset I \times J := \{1, \dots, M\} \times \{1, \dots, n\}$ be index sets such that

$$\forall (i, j) \in IJ_D : \forall \mathbf{x} \in \Gamma_i : \quad \mathbf{b}_i^{(j)}(\mathbf{x}) \cdot \mathbf{u}(\mathbf{x}) = g_i^{(j)}(\mathbf{x}), \quad (14)$$

$$\forall (i, j) \in IJ_R : \forall \mathbf{x} \in \Gamma_i : \quad \frac{\partial (\mathbf{b}_i^{(j)}(\mathbf{x}) \cdot \mathbf{u}(\mathbf{x}))}{\partial \mathbf{n}} + \mathbf{c}_i^{(j)}(\mathbf{x}) \cdot \mathbf{u}(\mathbf{x}) = h_i^{(j)}(\mathbf{x}), \quad (15)$$

Without loss of generality, we assume that for every $(i, j) \in IJ_R$ and $\mathbf{x} \in \Gamma_i$ the $c_i^{(j)}(\mathbf{x})$ lies in $\text{span}(\{\mathbf{b}_i^{(k)}(\mathbf{x}) \mid (i, k) \notin IJ_D\})$. The generic solution structure now becomes

$$\mathbf{u}(\mathbf{x}) = \sum_{i=1}^M w_i(\mathbf{x}) \mathbf{u}_i(\mathbf{x}) + [\Psi^{(1)}(\mathbf{x}), \dots, \Psi^{(n)}(\mathbf{x})]^T \prod_{i=1}^M \phi_i(\mathbf{x})^{\mu_i}, \quad (16)$$

$$\mu_i = \begin{cases} 2, & \exists j : (i, j) \in IJ_R, \\ 1, & \text{else,} \end{cases} \quad (17)$$

with weights given in (7). Functions $\mathbf{u}_i$ are expressed as a linear combination of the basis functions

$$\forall i \in I : \quad \mathbf{u}_i(\mathbf{x}) = \sum_{j=1}^n \mathbf{b}_i^{(j)}(\mathbf{x}) u_i^{(j)}(\mathbf{x}). \quad (18)$$

Proposition 2.2. If the vectors $\mathbf{b}_i^{(1)}(\mathbf{x}), \dots, \mathbf{b}_i^{(n)}(\mathbf{x})$ form a basis of $\mathbb{R}^n$ for every $\mathbf{x} \in N(\Gamma_i)$ where $N(\Gamma_i)$ is an open neighborhood of Γ_i , we have

$$\forall (i, j) \in I \times J : \quad \forall \mathbf{x} \in \Gamma_i : \quad \mathbf{b}_i^{(j)}(\mathbf{x}) \cdot \mathbf{u}(\mathbf{x}) = u_i^{(j)}(\mathbf{x}), \quad (19)$$

$$\forall (i, j) \in I \times J \text{ with } \phi_i^{\mu_i} \text{ having a vanishing gradient} : \quad \forall \mathbf{x} \in \Gamma_i : \quad \frac{\partial (\mathbf{b}_i^{(j)}(\mathbf{x}) \cdot \mathbf{u}(\mathbf{x}))}{\partial \mathbf{n}} = \frac{\partial u_i^{(j)}(\mathbf{x})}{\partial \mathbf{n}}. \quad (20)$$

Proof. The proof is shown in Appendix E. $\square$

Therefore, every $u_i^{(j)}$ has to satisfy the corresponding boundary condition. For Dirichlet conditions this is achieved by setting

$$\forall (i, j) \in IJ_D : \quad u_i^{(j)}(\mathbf{x}) = \begin{cases} g_i^{(j)}(\mathbf{x}) + \bar{\phi}_i(\mathbf{x}) \tilde{\Psi}_i^{(j)}(\mathbf{x}), & \phi_i^{\mu_i} \text{ has a vanishing gradient,} \\ g_i^{(j)}(\mathbf{x}), & \text{else,} \end{cases} \quad (21)$$

for both GLSS and OP. However, the two approaches differ in their treatment of Robin conditions.

2.2.1 GLSS

If all boundary segments $\Gamma_1, \dots, \Gamma_M$ are pairwise disjoint, we set

$$u_i^{(j)}(\mathbf{x}) = \Psi_i^{(j)}(\mathbf{x}) - \bar{\phi}_i(\mathbf{x}) \nabla \bar{\phi}_i(\mathbf{x}) \cdot \nabla \Psi_i^{(j)}(\mathbf{x}) + \bar{\phi}_i(\mathbf{x}) f_i^{(j)}(\mathbf{x}), \quad (22)$$

$$f_i^{(j)}(\mathbf{x}) = \mathbf{c}_i^{(j)}(\mathbf{x}) \cdot \sum_{k=1, (i,k) \notin IJ_D}^n \mathbf{b}_i^{(k)}(\mathbf{x}) \Psi_i^{(k)}(\mathbf{x}) - h_i^{(j)}(\mathbf{x}). \quad (23)$$

for $(i, j) \in IJ_R$. For $(i, j) \in I \times J \setminus (IJ_D \cup IJ_R)$, we define $u_i^{(j)}(\mathbf{x}) = \Psi_i^{(j)}(\mathbf{x})$. As above, the functions $\Psi_i^{(j)}$ have to be modified, if Γ_i has intersection points with other segments.

Intersecting boundary segments. We generalize our approach to the system case and let

$$\Psi_i^{(j)}(\mathbf{x}) = \mathbf{b}_i^{(j)}(\mathbf{x}) \cdot \left(\frac{\phi_B(\mathbf{x})}{\phi_A(\mathbf{x}) + \phi_B(\mathbf{x})} \mathbf{u}_A(\mathbf{x}) + \frac{\phi_A(\mathbf{x})}{\phi_A(\mathbf{x}) + \phi_B(\mathbf{x})} \mathbf{u}_B(\mathbf{x}) \right) + \phi_A(\mathbf{x}) \phi_B(\mathbf{x}) \bar{\Psi}_i^{(j)}(\mathbf{x}). \quad (24)$$

The construction of the functions $\mathbf{u}_A$ and $\mathbf{u}_B$ is more difficult than in the scalar case. We demonstrate how to do this with an example. Consider two segments Γ_1 and Γ_2 with intersection point P and assume for simplicity that $u \in \mathbb{R}^3$. Let Dirichlet conditions be prescribed on Γ_1 with respect to the basis vectors $\mathbf{b}_1^{(1)}(P) = (1, 0, 0)^T$ and $\mathbf{b}_1^{(2)}(P) = (0, 1, 0)^T$ and on Γ_2 with respect to the vector $\mathbf{b}_2^{(1)}(P) = (1, 1, 0)^T$. These three basis vectors span a two-dimensional subspace with basis $(1, 0, 0)^T, (0, 1, 0)^T$. We define $\mathbf{u}_P$ as a linear combination of this basis and an unknown component acting on the orthogonal complement, i.e.

$$\mathbf{u}_P(\mathbf{x}) = g_P^{(1)} \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} + g_P^{(2)} \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} + \Psi_P^{(3)}(\mathbf{x}) \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}. \quad (25)$$

Note that the constants $g_P^{(1)}$ and $g_P^{(2)}$ have to be chosen such that $\mathbf{u}_P$ satisfies all Dirichlet conditions prescribed in P . A complete algorithm can be found in the Appendix C.

Table 1: Size, training and inference times of the FNO for the four different approaches to enforce boundary conditions for the Darcy flow and Navier-Stokes equations.

	Trainable parameters	Checkpoint size (MByte)	Training time (min)	Inference time (sec)
Darcy flow				
GLSS	13,132,932	105	182.35	0.0130
OP	13,132,674	105	181.68	0.0120
Semi-weak	13,132,545	105	180.79	0.0104
Weak	13,132,545	105	181.44	0.0101
Navier-Stokes equations				
GLSS	13,133,706	105	9.79	0.0298
OP	13,133,190	105	9.76	0.0285
Semi-weak	13,132,803	105	9.49	0.0248
Weak	13,132,803	105	9.62	0.0238

2.2.2 OP

If all boundary segments Γ_i lie in hyperplanes and $\mathbf{b}^{(j)} := \mathbf{b}_1^{(j)} = \dots = \mathbf{b}_M^{(j)}$ holds for every $j = 1, \dots, n$, the global solution structure, given by (16) and (18), simplifies to

$$\mathbf{u}(\mathbf{x}) = \sum_{j=1}^n \mathbf{b}^{(j)}(\mathbf{x}) \sum_{i=1}^M w_i(\mathbf{x}) u_i^{(j)}(\mathbf{x}) + [\Psi^{(1)}(\mathbf{x}), \dots, \Psi^{(n)}(\mathbf{x})]^T \prod_{i=1}^M \phi_i(\mathbf{x})^{\mu_i}. \quad (26)$$

We choose the local solution structures as

$$u_i^{(j)}(\mathbf{x}) = \begin{cases} g_i^{(j)}(\mathbf{x}) + \bar{\phi}_i(\mathbf{x}) \tilde{\Psi}_i(\mathbf{x}), & \phi_i \text{ has a vanishing gradient} \\ g_i^{(j)}(\mathbf{x}), & \text{else} \end{cases}, \quad (i, j) \in IJ_D, \quad (27)$$

$$\begin{cases} \tilde{\Psi}^{(j)}(\mathcal{N}(\mathbf{x}; \bar{\phi}_i)) + \bar{\phi}_i(\mathbf{x}) f_i^{(j)}(\mathcal{N}(\mathbf{x}; \bar{\phi}_i)), & (i, j) \in IJ_R, \\ \tilde{\Psi}^{(j)}(\mathbf{x}), & \text{else} \end{cases}$$

The functions $f_i^{(j)}$ are defined as

$$f_i^{(j)}(\mathbf{x}) = \mathbf{c}_i^{(j)}(\mathbf{x}) \cdot \sum_{k=1, (i,k) \notin IJ_D}^n \tilde{\Psi}^{(k)}(\mathbf{x}) \mathbf{b}^{(k)}(\mathbf{x}) - h_i^{(j)}(\mathbf{x}), \quad (28)$$

and the $\tilde{\Psi}^{(j)}$ are defined as

$$\tilde{\Psi}^{(j)}(\mathbf{x}) = g^{(j)}(\mathbf{x}) + \tilde{\Psi}^{(j)}(\mathbf{x}) \prod_{i=1, (i,j) \in IJ_D}^M \phi_i(\mathbf{x}). \quad (29)$$

Each function $g^{(j)}$ is chosen in a way that it satisfies all Dirichlet conditions prescribed with respect to the basis vector $\mathbf{b}^{(j)}$ and $\tilde{\Psi}^{(j)}$ is an unknown function to be approximated.

Theorem 2.3. The derived solution structure satisfies the boundary conditions (14) and (15) for both the GLSS and OP approach.

Proof. The proof is in Appendix F. □

3 Numerical results

Details regarding the architecture and training of the network can be found in Appendix G. Table 1 shows different training-related parameters of the networks arising from the four approaches for the two benchmarks. Because the size of the network is dominated by the size of the four convolution layers, the number of trainable parameters varies only very slightly. There is no discernible impact on the size of the checkpoint files. Training times are stable, with semi-weak boundary conditions training the fastest in both cases but the difference to the slowest GLSS is below 3%. Inference times increase for GLSS and OP compared to weakly enforced boundary conditions. We see the largest increase by about 29% for GLSS for the Darcy flow. Note that training for the Navier-Stokes equations is much faster because we train only solutions and no solution operator.




Figure 3: Training progress and errors on the validation set for different ways to enforce boundary conditions for the Darcy flow.

3.1 Darcy Flow

As scalar test problem, we consider the Darcy flow equation governing fluid flow in porous media Darcy (1856). The detailed numerical setup can be found in Appendix H, including the precise local solution structures used for GLSS and OP. Figure 3 shows training loss (upper) and validation error (lower). The left column shows the loss and error curve for the PINN-like training. Both these curves are the average loss and l^2 -error over the 100 parameters that the PINO was trained on individually in PINN-style. The middle column shows the loss and error curve for PINO trained on 400 parameters. The right column shows the loss and error curve for finetuning. Dotted lines indicate the loss and error value at the very beginning of the finetuning.

All cases train reasonably well, reducing the loss function by at least one order of magnitude (OP for PINN-like training) and two or more orders in most cases. Losses are not indicative of achieved validation errors. For the PINN-style training, OP and GLSS are more accurate than weak or semi-weak boundary conditions. The same holds true for finetuning, where OP is slightly more accurate than GLSS. For operator training, OP is more accurate than GLSS which performs on par with semi-weak and better than weak boundary conditions. In summary, for the Darcy flow, even though losses do not necessarily decay faster, OP and GLSS in almost all cases produces more accurate solutions than weak or semi-weak boundary conditions. Table 2 shows the average l_2 -error plus standard deviation (left column), best case l_2 -error (middle column) and worst case l_2 -error (right column). For operator training and finetuning, OP is the most accurate approach whilst GLSS is the most accurate for PINN-like training. For best case errors, shown in the middle column, there is no clear benefit from the two new approaches. However, there are substantial gains in accuracy from OP and GLSS for the worst case in PINN-like training and finetuning and from OP in operator training. Plots of the median, best- and worst-case solutions can be found in Appendix I.

Table 2: l_2 -errors of the predicted u against the analytical solution for the four different approaches to enforce boundary conditions for the Darcy flow problem.

	Operator training			Finetuning			PINN-like training		
	Average	Best	Worst	Average	Best	Worst	Average	Best	Worst
GLSS	0.03±0.04	0.004	0.27	0.01±0.01	0.003	0.06	0.02±0.02	0.005	0.08
OP	0.02±0.01	0.003	0.06	0.01±0.01	0.002	0.04	0.04±0.02	0.006	0.11
S-Weak	0.03±0.03	0.002	0.17	0.03±0.02	0.003	0.10	0.09±0.03	0.021	0.17
Weak	0.05±0.05	0.008	0.28	0.04±0.05	0.003	0.26	0.13±0.14	0.008	0.64




Figure 4: Training loss (upper) and validation error in the u (upper middle) and v (lower middle) velocity component and pressure p (lower) for the Navier-Stokes equations.

3.2 Navier-Stokes equations

We use the standard benchmark by Turek, et al. (1996), simulating 2D stationary flow through a channel and across a cylinder. The details of the problem setup can be found in Appendix J, including the precise form of the local solution structures used for GLSS and OP.

Figure 4 shows the training losses in the upper figure and the errors in the velocity components u and v and the pressure p against the numerically computed reference solution. After 4000 epochs, losses for GLSS, OP and weak boundary conditions are similar but the loss for semi-weak remains somewhat higher. In terms of errors, we again see a clear benefit in terms of accuracy from GLSS and OP as they outperform weak and semi-weak boundary conditions in all three solution components.

To further assess accuracy we consider three practically relevant diagnostic quantities: pressure difference, drag coefficient and lift coefficient, see Turek, et al. (1996) for their definition. Table 3 shows the values computed from the PINO using the four different ways to enforce boundary conditions and, in brackets, the relative error against the reference values by Nabh (1998). We again see a noticeable increase in accuracy from GLSS and OP over weak or semi-weak boundary conditions. Pressure difference and drag coefficient are predicted with high accuracy. While relative errors for the lift coefficient are large, they are still orders of magnitude smaller than for the weak or semi-weak approach.

Table 3: Physically important parameters computed from the Navier-Stokes solution. The reference values are provided by Nabh (1998) with 9 digit accuracy and we rounded them to 4 digits. The relative error against those reference values is shown in brackets.

	Pressure difference	Drag coefficient	Lift coefficient
GLSS	0.1150 (2.1%)	5.5336 (0.8%)	-0.0058 (155%)
OP	0.1145 (2.6%)	5.5366 (0.8%)	0.0024 (77%)
Semi-weak	0.0678 (42.3%)	3.8221 (31.5%)	-0.3759 (3646%)
Weak	0.0902 (23.2%)	4.6633 (16.4%)	0.3849 (3531%)
Reference values	0.1175	5.5795	0.0106

References

- Alnaes, M. S. & Logg, A. & Ølgaard, K. B. & Rognes, M. E. & Wells, G. N. (2014) Unified Form Language: A domain-specific language for weak formulations of partial differential equations, *ACM Transactions on Mathematical Software* 40. [doi.org/10.1145/2566630].
- Baratta, I. A. & Dean, J. P. & Dokken, J. S. & Habera, M. & Hale, J. S. & Richardson, C. N. & Rognes, M. E. & Scroggs, M. W. & Sime, N. & Wells, G. N. (2023) DOLFINx: The next generation FEniCS problem solving environment, preprint. [doi.org/10.5281/zenodo.10447666].
- Berrone, S. & Canuto, C. & Pintore, M. & Sukumar, N. (2023). Enforcing Dirichlet boundary conditions in physics-informed neural networks and variational physics-informed neural networks. <https://doi.org/10.1016/j.heliyon.2023.e18820>.
- Biswas, A. & Shapiro, V. (2004). Approximate distance fields with non-vanishing gradients. <https://doi.org/10.1016/j.gmod.2004.01.003>.
- Gladstone, R.J. & Nabian, M.A. & Sukumar, N. & Srivastava, A. & Meidani, H. (2022). FO-PINNs: A First-Order formulation for Physics Informed Neural Networks. *ArXiv*, abs/2210.14320.
- Kingma, D. P. & Ba, J. (2017). Adam: A Method for Stochastic Optimization. <https://arxiv.org/abs/1412.6980>.
- Li, Z. & Kovachki, N.B. & Azizzadenesheli, K. & Liu, B. & Bhattacharya, K. & Stuart, A.M. & Anandkumar, A. (2020a) Neural Operator: Graph Kernel Network for Partial Differential Equations. *ArXiv*, abs/2003.03485.
- Li, Z. & Kovachki, N.B. & Azizzadenesheli, K. & Liu, B. & Bhattacharya, K. & Stuart, A.M. & Anandkumar, A. (2020b) Fourier Neural Operator for Parametric Partial Differential Equations. *ArXiv*, abs/2010.08895.
- Li, Z. & Zheng, H. & Kovachki, N.B. & Jin, D. & Chen, H. & Liu, B. & Azizzadenesheli, K. & Anandkumar, A. (2021) Physics-Informed Neural Operator for Learning Partial Differential Equations. *ArXiv*, abs/2111.03794.
- Li, Z. & Huang, D.Z. & Liu, B. & Anandkumar, A. (2022) Fourier Neural Operator with Learned Deformations for PDEs on General Geometries. *J. Mach. Learn. Res.*, 24, 388:1-388:26.
- Li, Z. & Kovachki, N. B. & Choy, C. & Li, B. & Kossai, J. & Otta, S. P. & Nabian, M. A. & Stadler, M. & Hundt, C. & Azizzadenesheli, K. & Anandkumar, A. (2023). Geometry-Informed Neural Operator for Large-Scale 3D PDEs. <https://arxiv.org/abs/2309.00583>.
- Lu, L. & Jin, P. & Pang, G. & Zhang, Z. & Karniadakis, G. E. (2021). Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators. <https://doi.org/10.1038/s42256-021-00302-5>
- Lu, L. & Pestourie, R. & Yao, W. & Wang, Z. & Verdugo, F. & Johnson, S. G. (2021). Physics-Informed Neural Networks with Hard Constraints for Inverse Design. <https://doi.org/10.1137/21M1397908>.
- Lu, L. & Meng, X. & Cai, S. & Mao, Z. & Goswami, S. & Zhang, Z. & George Em Karniadakis, G.E. (2022) A comprehensive and fair comparison of two neural operators (with practical extensions) based on FAIR data. *ArXiv*. abs/2111.05512.
- Manavi, S. & Fattahi, E. & Becker, T. (2024). A trial solution for imposing boundary conditions of partial differential equations in physics-informed neural networks. <https://doi.org/10.1016/j.engappai.2023.107236>.
- McFall, K. S. & Mahan, J. R. (2009). Artificial Neural Network Method for Solution of Boundary Value Problems With Exact Satisfaction of Arbitrary Boundary Conditions. <https://doi.org/10.1109/TNN.2009.2020735>.
- Nabh, G. (1998) On High Order Methods for the Stationary Incompressible Navier-Stokes Equations; University of Heidelberg; Preprint 42/98.

310 Raissi, M. & Perdikaris, P. & Karniadakis, G. E. (2019). Physics-informed neural networks: A
311 deep learning framework for solving forward and inverse problems involving nonlinear partial
312 differential equations. <https://doi.org/10.1016/j.jcp.2018.10.045>.

313 Raissi, M. & Perdikaris, P. & Ahmadi, N. & Karniadakis, G. E. (2024). Physics-Informed Neural
314 Networks and Extensions. <https://arxiv.org/abs/2408.16806>.

315 Rvachev, V. L. & Sheiko, T. I. (1995) R-Functions in Boundary Value Problems in Mechanics. ASME.
316 Appl. Mech. Rev. April 1995; 48(4): 151–188. <https://doi.org/10.1115/1.3005099>.

317 Rvachev, V. & Sheiko, T. & Shapiro, V. & Tsukanov, Igor. (2000) On completeness of RFM solution
318 structures. Computational Mechanics. 25. 305-317. 10.1007/s004660050479.

319 Rvachev, V. L. & Sheiko, T. I. & Shapiro, V. & Tsukanov, I. (2001) Transfinite interpolation
320 over implicitly defined sets. Computer Aided Geometric Design; Volume 18, Issue 3: 195-220.
321 [https://doi.org/10.1016/S0167-8396\(01\)00015-2](https://doi.org/10.1016/S0167-8396(01)00015-2).

322 Scroggs, M. W. & Baratta, I. A. & Richardson, C. N. & Wells, G. N. (2022) Basix: a
323 runtime finite element basis evaluation library, Journal of Open Source Software 7(73) 3982.
324 [doi.org/10.21105/joss.03982].

325 Scroggs, M. W. & Dokken, J. S. & Richardson, C. N. & Wells, G. N. (2022) Construction of
326 arbitrary order finite element degree-of-freedom maps on polygonal and polyhedral cell meshes,
327 ACM Transactions on Mathematical Software 48(2). 18:1–18:23. [doi.org/10.1145/3524456].

328 Straub, C. & Brendel, P. & Medvedev, V. & Roskopf, A. (2025). Hard-constraining Neumann
329 boundary conditions in physics-informed neural networks via Fourier feature embeddings.
330 <https://arxiv.org/abs/2504.01093>.

331 Sukumar, N. & Srivastava, A. (2021) Exact imposition of boundary conditions with distance functions
332 in physics-informed deep neural networks. ArXiv, [abs/2104.08426](https://arxiv.org/abs/2104.08426).

333 Toscano, J. D. & Oommen, V. & Varghese, A. J. & Zou, Z. & Ahmadi Daryakenari, N. & Wu, C. &
334 Karniadakis, G. E. (2025). From PINNs to PIKANs: recent advances in physics-informed machine
335 learning. <https://doi.org/10.1007/s44379-025-00015-1>.

336 Turek, S. & Schaefer, M. (1996) Benchmark computations of laminar flow around cylinder; in
337 Flow Simulation with High-Performance Computers II, Notes on Numerical Fluid Mechanics 52,
338 547-566, Vieweg 1996.

339 Zeinhofer, M. & Masri, R. & Mardal, K.-A. (2024). A Unified Framework for the Error Analysis of
340 Physics-Informed Neural Networks. <https://arxiv.org/abs/2311.00529>.

341 Darcy, H. (1856). Les fontaines publiques de la ville de Dijon: exposition et application des principes
342 à suivre et des formules à employer dans les questions de distribution d’eau.

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: The central claim of our paper is that the two new proposed approaches can overcome the stability issues for non- C^1 boundaries that previous ways to prescribe boundary conditions suffer from. We provide numerical evidence for two challenging benchmarks that this is the case.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Limitations of the two new approaches are concisely summarized at the end of the introduction, right after stating the contributions of the paper.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: All propositions are either proved in the paper or we cite proofs from the literature.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Yes, we strive to provide all required parameters to reproduce the studied setups in either the paper or the supplementary appendix. Should it turn out that we missed some parameter, the code is made available for reference.

Guidelines: All experimental results in the paper come from numerical experiments and can be reproduced using the provided code, see the answer to the next question.

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.

sions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example

- (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
- (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The reviewers can access all the code required to reproduce the presented results via FigShare: <https://figshare.com/s/6332d1c1e782304fb264>. In case of publication, the code will be provided via GitHub and a Zenodo-provided DOI and cited in the camera-ready version of the paper. The trained model parameters will be made freely available via an institutional repository and be cross-linked with the paper and the code.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- [public/guides/CodeSubmissionPolicy](https://nips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to [//nips.cc/public/guides/CodeSubmissionPolicy](https://nips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: We carefully included all the settings used for training so that readers can understand and reproduce the results. In case we missed some crucial piece of information, the full code is available for reference.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [NA] .

Justification: The results in the paper are not statistical in nature and we do not rely on statistical significant to evidence our claims. While we show solutions for a randomly generated set of parameters, we present and discuss average results as well as best and worst case results, thus providing a complete picture of the spread of results obtained by our method.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Training and inference times as well as the number of trainable parameters of the used networks is stated in the paper. Details on the hardware used for training (a standard commodity desktop PC) are provided in the Appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the

Answer: [Yes]

Justification: Our research does not involve human subjects or uses any data that could raise privacy concerns. Given its fairly theoretical and fundamental nature, we also cannot envision any scenario in which it could cause societal harm.

Guidelines:

- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: Our research is foundational and addresses the solution of partial differential equations via machine learning. While we use two benchmarks from computational fluid dynamics to assess performance, these are not tied to any specific application. We therefore cannot see any direct pathways how our research could lead to any form of harm.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: Because we do not expect any harmful impact of our research, we have not put any safeguards in place but opted to publish the fully trained models together with all the code necessary to reproduce our experiments.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.

- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Our code uses three other software packages, pyTorch (<https://github.com/pytorch/pytorch>), the FNO code by Li et al. (https://github.com/neuraloperator/physics_informed) and FEnICSx (<https://docs.fenicsproject.org/>). Our usage complies with their licenses: a bespoke license for pyTorch Apache-2 for FNO and MIT license for FEnICSx.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The provided code comes with in-code comments as well as a README that provides guidance how to reproduce the results shown in the paper.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: Neither crowdsourcing nor human subjects played any role in this research.

Guidelines:

- 655 • The answer NA means that the paper does not involve crowdsourcing nor research with
656 human subjects.
657 • Including this information in the supplemental material is fine, but if the main contribu-
658 tion of the paper involves human subjects, then as much detail as possible should be
659 included in the main paper.

661 or other labor should be paid at least the minimum wage in the country of the data
662 collector.

663 **15. Institutional review board (IRB) approvals or equivalent for research with human**
664 **subjects**

665 Question: Does the paper describe potential risks incurred by study participants, whether
666 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
667 approvals (or an equivalent approval/review based on the requirements of your country or
668 institution) were obtained?

669 Answer: [NA]

670 Justification: Neither crowdsourcing nor human subjects played any role in this research.

671 Guidelines:

- 672 • The answer NA means that the paper does not involve crowdsourcing nor research with
673 human subjects.
674 • Depending on the country in which research is conducted, IRB approval (or equivalent)
675 may be required for any human subjects research. If you obtained IRB approval, you
676 should clearly state this in the paper.
677 • We recognize that the procedures for this may vary significantly between institutions

679 guidelines for their institution.
680 • For initial submissions, do not include any information that would break anonymity (if
681 applicable), such as the institution conducting the review.

682 **16. Declaration of LLM usage**

683 Question: Does the paper describe the usage of LLMs if it is an important, original, or
684 non-standard component of the core methods in this research? Note that if the LLM is used
685 only for writing, editing, or formatting purposes and does not impact the core methodology,
686 scientific rigorousness, or originality of the research, declaration is not required.

687 Answer: [NA]

688 Justification: LLMs played no role in the development of the presented methodology.

689 Guidelines:

- 690 • The answer NA means that the core method development in this research does not
691 involve LLMs as any important, original, or non-standard components.

693 for what should or should not be described.