
FLEXOLMO: Open Language Models for Flexible Data Use

Abstract

We introduce **FLEXOLMO**, a new class of language models (LMs) that supports (1) *distributed training without data sharing*, where different model parameters are independently trained on closed datasets, and (2) *data-flexible inference*, where these parameters along with their associated data can be flexibly included or excluded from model inferences with no further training. FLEXOLMO employs a mixture-of-experts (MoE) architecture where each expert is trained independently on closed datasets and later integrated through a new domain-informed routing without any joint training. FLEXOLMO is trained on FLEXMIX, a corpus we curate comprising publicly available datasets alongside seven domain-specific sets, representing realistic approximations of closed sets. We evaluate models with up to 37 billion parameters (20 billion active) on 31 diverse downstream tasks. We show that a general expert trained on public data can be effectively combined with independently trained experts from other data owners, leading to an average 41% relative improvement while allowing users to opt out of certain data based on data licensing or permission requirements. Our approach also outperforms prior model merging methods by 10.1% on average and surpasses the standard MoE trained without data restrictions using the same training FLOPs. Altogether, this research presents a solution for both data owners and researchers in regulated industries with sensitive or protected data. FLEXOLMO enables benefiting from closed data while respecting data owners’ preferences by keeping their data local and supporting fine-grained control of data access during inference.

1 Introduction

Pretraining language models (LMs) typically requires centralized access to all data during training and does not have any mechanism to track or control the influence of specific data points on model parameters. Model developers must therefore make a one-time decision on which data sources to

*Core contributors.




Figure 1: **An overview of FLEXOLMO.** Data owners can contribute without sharing the data by training their own expert modules (FFNs and router embeddings) with a shared public model as an anchor point. At inference, these modules are integrated into a MoE model via a novel router embedding concatenation. This design enables flexible inclusion or exclusion of experts and strict opt-out guarantees, e.g., Github data can be excluded at no cost (blurred) during inference.

include, with limited ability to remove the effect of certain data after training [1, 2, 3]. Moreover, this centralized approach precludes the use of closed data that data owners cannot share with model developers for confidentiality, regulatory, or other reasons. Although solutions have been proposed to allow training without sharing the data, such as federated learning [4, 5], their practical adoption remains limited due to performance degradation and the high cost of synchronized training [6, 7].

We introduce FLEXOLMO, a new class of LMs that enables distributed training on locally maintained datasets while enabling flexible opt-in and opt-out during inference. FLEXOLMO (Figure 1) employs a mixture-of-experts (MoE) architecture [8, 9], where each expert is trained independently on closed datasets and later integrated into an MoE. This design allows data owners to contribute asynchronously without sharing their data, while also enabling continual updates with new data and providing strong guarantees for data opt-out during inference. Our approach can be seen as an instance of model merging [10], which merges different models into a unified one [11, 12]. However, our model is designed to address the unique challenges in our problem setup—combining models pre-trained on completely disjoint datasets with different distributions—which makes prior model merging techniques like ensembling output probabilities [11] or merging model weights [12] suboptimal.

A key challenge in training FLEXOLMO is ensuring the merging of independently trained experts without joint training. We introduce a training algorithm where each data owner independently trains an expert module using the frozen public model as a shared anchor (Figure 1). This approach teaches independently trained experts to coordinate with the same public model and, by extension, with each other. Additionally, the router, a module that determines which experts process each token, typically requires joint training. We address this by assigning each expert a router embedding, initialized from its domain embedding using an off-the-shelf embedder [13] and further finetuned on its corresponding data during individual expert training. These embeddings are then concatenated to form the router during merging, removing the need for joint training.

To validate FLEXOLMO, we curate a data mixture called FLEXMIX, which includes a public training set along with seven domain-specific sets (e.g., news, educational text, and Reddit). These domains are chosen to simulate scenarios where high-quality data that can benefit LM training is not publicly available.

We train FLEXOLMO first on public data, then extend it by merging expert modules trained independently on our simulated closed sets. While continued pretraining on these sets improves some downstream tasks, it suffers from catastrophic forgetting and inconsistent performance. In contrast, FLEXOLMO improves upon the public model by 41% and also outperforms prior merging techniques such as model soup and ensembling by 10.1% across 31 downstream tasks. We observe the largest improvements on tasks related to closed sets. Notably, even on benchmarks where no individual

closed set improved performance over the public model, combining multiple experts yielded significant gains, demonstrating synergies among independently trained modules. Our qualitative analysis demonstrates that sparse expert activation across layers through the MoE architecture is key to these gains. Our qualitative analysis shows that the MoE architecture’s ability to selectively activate different experts per layer per token is crucial to these gains by combining the strengths of each specialized expert. We hope our work enables research with a broader range of closed datasets for LM training, particularly for organizations interested in collaborating on scientific research through the new features that FlexOlmo provides.

2 Background & Related Work

2.1 Background: Data Restrictions

The standard LM training practice requires model developers to aggregate all data centrally and make a one-time decision on which data source to include and exclude. But many real-world data come with sharing and usage restrictions and necessitates (1) model training without data pooling and (2) model inference that can flexibly select different data sources based on use case and access privileges.

Data Sharing Constraints Organizations in regulated industries require LMs that can leverage their closed datasets while maintaining strict data privacy and access controls. Healthcare institutions, financial firms, and other entities possess valuable domain-specific data but cannot share it externally due to HIPAA, GDPR [14, 15], data sovereignty laws [16], and intellectual property (IP) protections. These organizations need training paradigms that enable AI improvement on their sensitive data while ensuring such sensitive data never leaves certain environments and can be removed from the model after training, e.g., when data usage rights expire. In such settings, modular training approaches, where individual experts are trained independently and asynchronously on locally maintained data, are essential.

Data Use Constraints The inclusion of certain data depends on specific use cases and end users. *Privileged access:* User-facing applications often involve closed data restricted to specific, authorized users [17]. For example, GitHub Copilot must tailor code suggestions to reflect internal repositories based on an engineer’s role and access rights [18]. *Copyright and data consent:* Legal and ethical considerations on training data for AI are evolving and uncertain [19, 20, 21, 22, 23, 24], and often depend on the data’s intended use, e.g., licenses may prohibit commercial use or limit certain query types [25, 26]. *Model control:* Training data often include sensitive content [27, 28, 29] which may be beneficial in certain contexts but harmful in others. For instance, one may want to activate the use of toxic content for toxicity classification in a research setting, but deactivate it in applications presented to a general audience.

2.2 Related Work

Federated Learning Federated Learning (FL) trains a single model over distributed datasets by synchronously aggregating client updates [5, 4, 30]. FL methods range from classical approaches that iteratively aggregate parameter updates from local clients [5, 31] to parameter-efficient techniques which have been adapted for LMs [32, 33, 34]. FL can guarantee data privacy using techniques such as homomorphic encryption [35] and differential privacy (DP) [36]. However, FL has seen limited adoption in LM training due to the high cost of synchronization and performance degradation [6, 7], and remains susceptible to privacy attacks due to inter-client communication [37, 38].

Similar to [39], our approach avoids data sharing but differs fundamentally by supporting independent, asynchronous training without costly inter-client communication, and allowing real-time opt-in and opt-out. Like FL, our model allows data owners to optionally apply DP training locally for privacy guarantees. Because DP is orthogonal to our architecture, each contributor can independently choose whether to apply it, providing flexibility without compromising the overall design.

Model Merging Our work builds on recent efforts [40, 10] that advocate for developing machine learning models like open-source softwares, where sub-parts of the model can be trained independently and subsequently merged into unified systems. This can be achieved through various methods, including weight merging, output ensembling, and expert routing. Model soup—merging model weights trained on different datasets from the same initialization—can boost performance [41, 12,

42, 43, 44], especially with weighted combinations [45, 46, 47, 48, 49, 50, 51]. Weighted output ensembling (e.g., BTM [52, 53]) is also effective when models are trained on distinct datasets initialized from the same seed model. These approaches can be applied to our setting, where each expert is independently trained starting from the same public model then merged into a unified one. Our experiments (§5) show that these methods are less effective, primarily because they lack learned connections between different modules, which constrains the expressivity of the resulting models.

An alternative line of work focuses on expert routing methods, such as DEMix [54], BTX [55] and its extensions [56, 57, 58], which merge dense, independently trained experts into a mixture-of-experts (MoE) framework. We draw inspiration from this work, as we also integrate independently trained models into a MoE. However, these methods require joint training on a union of all datasets used in expert training after merging. By contrast, FLEXOLMO removes the need for joint data access to enable training on locally maintained datasets. Our work is also related to ModuleFormer [59], which induces sparse modularity from uncured data using novel load balancing and concentration losses.

Related efforts in parameter-efficient training have explored merging LoRA adapter weights trained on separate datasets [60, 61, 62, 63], particularly to reduce communication overhead in collaborative settings and support fine-grained data access control and opt-out use cases [64, 65]. Unlike these methods, which focus on merging lightweight adapters, our approach merges full expert models into a standard MoE architecture.

Mixture-of-Experts (MoE) MoE models [8, 9, 66], consisting of many small feedforward networks called *experts*, have gained popularity for their training and inference efficiency. Our work leverages the MoE architecture; however, our motivation and training method are fundamentally different as our primary goal is to support modularity rather than efficiency.

3 FLEXOLMO: LMs with Flexible Data Use

3.1 Problem Setup

Let M_{pub} be a model trained on a publicly available dataset D_{pub} , and $\mathcal{D} = \{D_1, D_2, \dots, D_n\}$ represent a collection of locally maintained datasets with separate owners. Our objective is a single model M_{final} , which is constructed via composing M_{pub} and a set of modules $\{M_1, M_2, \dots, M_n\}$, where each M_i is independently trained by the owner of D_i , who also has access to M_{pub} .

This model satisfies two requirements: (1) training M_{final} does not require anyone to have joint access to the full dataset collection $\mathcal{D}$, as each M_i is trained independently by the owner of dataset D_i ; (2) removing any module M_i from M_{final} guarantees complete removal of its associated data D_i .

The key modeling challenges are: (1) to develop an algorithm that creates M_i using D_i and M_{pub} , and (2) to design the merging algorithm that combines $M_{\text{pub}}, M_1, M_2, \dots, M_n$ into M_{final} .

3.2 Model Architecture

FLEXOLMO follows the standard MoE architecture: it replaces the feedforward network (FFN) in each transformer block with a router and n small FFNs called expert modules $\{M_{\text{pub}}, M_1, \dots, M_n\}$. Note that we omit the layer index for each expert in our notation for simplicity. Given a processed input token embedding $\mathbf{x} \in \mathbb{R}^h$, the MoE module computes output representation $\mathbf{y}$:

$$\mathbf{y} = \sum_{i \in \text{Top}k(r(\mathbf{x}))} \text{softmax}(r(\mathbf{x})_i) M_i(\mathbf{x}),$$

where the router function r computes the expert probabilities from $\mathbf{x}$. Unlike standard MoEs where experts are trained jointly, our experts are trained asynchronously on distinct datasets $\{D_1, \dots, D_n\}$.

3.3 Training Algorithm

Standard MoEs train all experts and the router jointly on all data. In contrast, FLEXOLMO trains experts independently by teaching them to coordinate (§3.3.1) and merges them at inference using a domain-informed router (§3.3.2). Optional router tuning can further improve performance (§3.3.3).

3.3.1 Training Experts to Coordinate

A straightforward way to train each expert would be to directly continue to train each expert M_i on its own data D_i [52]. We found that this method causes the experts to diverge too much from one another and from the original seed model, which makes merging after isolated training difficult.

To prevent such divergence, we train experts independently while teaching them to coordinate (Figure 1). We use M_{pub} as an anchor that teaches experts to coordinate with M_{pub} and, by extension, with each other. Specifically, during training, for dataset D_i , we construct a MoE model with two expert modules—both initialized from the same FFNs from M_{pub} . During training, we freeze M_{pub} expert and the shared attention layer, while the other expert (M_i) is trained on D_i . As each data owner updates only their own FFNs while keeping all other parameters (those inherited from M_{pub} such as attention layer) frozen, the learned FFNs are designed to naturally coordinate with each other later during merging at inference time. Importantly, with this approach, a router is learned so that each expert can be integrated into a MoE architecture without additional training (details in §3.3.2).

3.3.2 Domain-Informed Router

The router plays a critical role in MoE: the router function r maps an input vector $\mathbf{x}$ to a distribution over expert modules, including the public model as one of the experts:

$$r(\mathbf{x}) = \mathbf{W}_r \mathbf{x}, \quad \mathbf{W}_r \in \mathbb{R}^{(n+1) \times h}$$

In typical MoEs, $\mathbf{W}_r$ is trained end-to-end alongside all expert modules, using access to the full training dataset. Instead, we decompose $\mathbf{W}_r$ into individual expert-specific router embeddings, where each row $\mathbf{r}_i$ represents the router embedding for expert M_i , learned only from D_i :

$$\mathbf{W}_r = \begin{bmatrix} \text{---} & \mathbf{r}_{\text{pub}} & \text{---} \\ \text{---} & \mathbf{r}_1 & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{r}_n & \text{---} \end{bmatrix}, \text{ where } \mathbf{r}_i = \frac{1}{|S_i|} \sum_{d_k \in S_i} \mathbf{E}(d_k) \in \mathbb{R}^h, S_i \subset D_i.$$

These router embeddings can be initialized by averaging domain-specific embeddings of samples from each D_i , obtained by encoding subsets of data using an off-the-shelf embedder $\mathbf{E}$ [67] that maps a document into an h -dimensional vector. This method is motivated by prior model merging work that leverages domain embeddings for routing [68, 42, 69, 70, 71, 72, 73].

During coordinated training of experts (§3.3.1), we learn the router embeddings in pairs: $[\mathbf{r}_{\text{pub}}, \mathbf{r}_i]$. The public embedding $\mathbf{r}_{\text{pub}}$ remains frozen across all experts, while $\mathbf{r}_i$ is finetuned separately alongside the parameters of M_i . At inference time, merging the expert embeddings into the complete router matrix $\mathbf{W}_r$ directly integrates all expert modules into one unified MoE. Furthermore, experts can be flexibly added or removed by simply adding or removing their corresponding router embedding.

Adding a Bias Term Unlike standard router learning that is learned among all experts jointly, coordinated training of experts only learns pairwise routing decisions between one expert and the public model. This means the model never directly compares experts M_1 and M_2 during training, potentially limiting generalization during inference. To alleviate this issue, we add a negative bias term b_i for each independent trained expert $\{M_1, M_2, \dots, M_n\}$. We select expert M_i when:

$$\mathbf{r}_i \cdot \mathbf{x} + b_i > \mathbf{r}_{\text{pub}} \cdot \mathbf{x} \quad \forall i \in \{1, 2, \dots, n\}$$

Otherwise default to M_{pub} . This helps the later merging process, where each expert competes not just with the public model but with all other experts. Further details and justifications are provided in §D.

3.3.3 Optional Router Training on Proxy Data

With our proposed model design, expert modules can be merged without any additional training. However, if data owners are willing to identify proxy samples within the public dataset M_{pub} that *resemble* their closed data, we can optionally perform a lightweight router tuning step after merging, using only public data from D_{pub} . Specifically, we assume each data owner selects a small proxy set $\hat{D}_i \subseteq D_{\text{pub}}$, where $|\hat{D}_i| \ll 0.01 \times |D_i|$, chosen to approximate the distribution of their closed dataset D_i . While $\hat{D}_i$ is too small to train expert modules, it still provides useful signals for improving router quality. To construct $\hat{D}_i$, we train a binary classifier to distinguish D_i from D_{pub} and select public samples with the highest predicted likelihood of belonging to D_i . After merging, we tune the router embeddings $\mathbf{r}_1, \dots, \mathbf{r}_n, \mathbf{r}_{\text{pub}}$ on the combined set $\hat{D}_1, \dots, \hat{D}_n$, and D_{pub} , sampled uniformly.

4 Experimental Setup

4.1 Training Data: FLEXMIX

Our corpus comprises a single *Public Mix* and seven *closed sets*, either real or simulated, which are designed to be disjoint from each other. Figure 5 in §B provides the statistics.

- **Public Mix** represents general web text based on Common Crawl (CC) ¹. Specifically, we took the Baseline version of DCLM [74], excluding news and creative writing content (described below). This represents a public dataset that can be used without restrictions.
- **News** includes news content from DCLM-Baseline, obtained by applying the classifier from [75] and selecting documents classified as News Articles. While included in CC when downloaded, many of the original sources are subject to closed access [20].
- **Creative Writing** includes creative content from DCLM-Baseline, obtained by applying the classifier from [75] and selecting documents classified as Creative Writing.
- **Code** includes code repositories from Starcoder [76, 77] with additional quality filtering as in [78].
- **Academic** includes open-access academic papers obtained from [79]; these are papers from [80, 81] but re-processed using olmOCR [79] for cleaner plain text.
- **Educational Text** includes educational text from digitized PDFs, converted to plain text using olmOCR [79].
- **Math** includes math-related content, including web pages about or using math and math problem sets, obtained by combining Dolma Math Mix [78] and FineMath4+ [82].
- **Reddit** contains posts and comments originally sourced and released by Dolma [83], further filtered and processed to improve quality (details in Appendix B). As of this writing, this Reddit data is no longer unrestrictedly downloadable due to Reddit’s 2023 policy change.²

These seven sets are designed to represent datasets with at least one of the following characteristics: (1) historically closed and not publicly available; (2) previously publicly available but now closed; or (3) domains with scarce high-quality public data.

4.2 Evaluation

We evaluate our models and baselines on a large and diverse collection of well-established benchmarks, consisting of 31 tasks across 10 categories, broadly grouped into (1) general-purpose LM benchmarks and (2) domain-specific evaluations. More details are provided in §C.

General-purpose Evaluation We report results on (1) **MC9**, nine multiple-choice datasets including ARC-Easy [84], ARC-Challenge [84], BoolQ [85], CSQA [86], HellaSwag [87], OpenBookQA [88], PIQA [89], SocQua [90], and WinoGrande [91], (2) **GEN5**, five generative tasks including CoQA [92], SQuAD [93], Natural Questions [94], TriviaQA [95], and DROP [96], as well as (3) **MMLU** [97], (4) **MMLU-Pro** [98], (5) **AGIEval** [99] consisting of 20 tasks from college admission tasks, and (6) **BBH** [100] consisting of 23 challenging BIG-Bench tasks.

Domain-specific Evaluation While general-purpose evaluation benchmarks already include some math assessment, we further evaluate math ability using (6) **Math2**, which encompasses two specialized math benchmarks: GSM8K [101] and MATH [102]. To evaluate coding capabilities, we use (7) **Code4**, 4 coding benchmarks including MBPP [103], MBPPPLUS [104], HUMAN-EVAL [105], and HUMAN-EVALPLUS [104]. To measure scientific literature understanding, we report on (8) **SciRIFF5**: comprising 5 subtasks from SciRIFF [106]. Finally, we include (9) **NewsG**: news generation and (10) **PoemG**: poem generation tasks, both evaluated using an LM judge.

4.3 Baselines

We compare our method against several baselines, either taken directly from prior work or initially adapted to our problem setting. All baselines, except for ‘Unrestricted training,’ train a set of dense

¹<https://commoncrawl.org>

²Reddit’s 2023 policy change restricts third-party access and use of its data, including for language model development; see nytimes.com/2023/04/18/technology/reddit-ai-openai-google.html.

models independently by continuing pretraining from the public model M_{pub} on each simulated closed set, without architectural changes, and merge them using model merging techniques.

Prompt-based Routing We use an LM-based domain classifier via prompting to route each query to the most suitable model, which is then used exclusively. We use Llama-3.1-8B-Instruct [107] and OLMo-2-1124-7B-Instruct [78] as classifiers. More details are in §A.1.

Model Soup We apply both average and weighted parameter averaging across all models, following [12]. The weights are derived by applying a softmax over the log-likelihoods of each model on the test example input.

Branch-Train-Merge (BTM) We follow BTM [11], which ensembles models by computing a weighted average of their output probabilities. Weights are obtained via a softmax over the log-likelihoods of the test example input of each model. As in the original BTM, ensembling can be restricted to the top- k models by zeroing the weights of all other models and renormalizing. See §A.1 for full details.

BTX We follow BTX [55], which upcycles an MoE from independently trained dense models. It copies the dense model parameters to the corresponding experts in MoE while averaging non-expert parameters such as attention layers for merging. The original BTX requires training all model parameters on combined datasets after merging. To approximate it as closely as possible while adhering to our setting, we perform this post-merge training on the public set only.

Unrestricted MoE To assess how closely our method approaches the benefits of full data access while preserving data separation, we construct an upper-bound reference model: a sparse MoE initialized from the public-only dense model and trained on the combined dataset, including all closed sets and Public Mix. As MoE training incurs roughly $2\times$ the FLOPs of our approach for the same data size, we report both compute-controlled ($1\times$ FLOPs, $0.5\times$ data) and data-controlled ($2\times$ FLOPs, $1\times$ data) comparisons.

4.4 Training Setup

For the public model M_{pub} , we use a dense model with 7 billion parameters following the OLMo 2 architecture [78]. This model contains 32 layers with hidden dimension 4,096 and is trained on our public mix for 1 trillion tokens. Following [78], we use a learning rate of 0.0009 and the AdamW optimizer with parameters $\beta_1 = 0.9$ and $\beta_2 = 0.95$ and a cosine learning rate scheduler. The public model is pretrained using 512 H100 GPUs with a global batch size of 4 million tokens for three days.

Each data owner then takes this checkpoint and performs continued-pretraining for 50 billion tokens on their own data (totaling 400B tokens across all experts). For the optional router training, we use 5 billion tokens in total. The final FLEXOLMO, trained on 8 sets, has 37 billion total parameters with 20 billion active (4 active experts out of 8). More details can be found in §A.2.

5 Results and Analysis

We conduct ablation studies and compare against a comprehensive set of baselines at a small scale with four experts—Public mix, math, educational text, and code (Table 1). We then evaluate our final model on the full setup including the Public mix and all seven simulated closed sets (Table 2). Finally, we present an in-depth analysis to illustrate the behavior and effectiveness of FLEXOLMO.

5.1 Main Results

Individual experts excel at their specialized tasks As shown in Table 1, experts trained on each domain-specific set demonstrate strong performance in their specialized domains: the Math expert achieves the highest scores on math tasks, while the Code expert performs the best on coding benchmarks. However, these experts exhibit considerable performance degradation when evaluated on tasks outside their domains. Notably, the Code expert performs poorly on general benchmarks.

FLEXOLMO outperforms individual experts FLEXOLMO outperforms individual experts in most cases. It improves upon the model trained solely on public data, achieving an average 41% relative gain. Largest improvements appear on benchmarks where closed data significantly boosts

Table 1: Evaluation of FLEXOLMO trained on *four* sets (public mix, math, educational text and code), tested on 24 tasks with 100 samples per subtask.

	MC9	GEN5	MMLU	MMLU Pro	AGI Eval	BBH	Math2	Code4	Avg.
Prev. Public model	68.4	58.8	57.0	27.1	39.0	35.6	8.1	1.0	36.9
<i>Individual experts</i>									
Math	63.8	46.3	51.1	24.0	40.7	45.4	50.4	18.1	42.5
Code	38.7	41.4	30.0	14.6	29.0	38.2	6.0	22.4	27.5
Educational Text	63.0	52.8	57.7	26.8	39.6	40.0	13.1	4.3	37.2
<i>Prior model merging work</i>									
Model soup (average)	70.6	53.8	54.7	28.4	41.4	42.4	17.5	8.2	39.6
Model soup (weighted)	69.6	56.3	58.6	30.5	45.4	43.5	18.5	14.8	42.2
BTM	69.0	58.5	59.6	29.0	43.6	43.6	21.2	22.3	43.4
Prompt-based routing (router: OLMo)	59.9	48.8	50.0	25.4	38.7	41.3	41.5	20.7	40.8
Prompt-based routing (router: Llama)	64.2	53.4	57.7	26.4	39.9	39.8	21.5	17.3	40.0
BTX	69.6	57.9	56.2	28.5	43.1	41.3	16.8	6.4	40.0
<i>Ours</i>									
FLEXOLMO (no optional router training)	71.1	58.6	58.1	28.4	44.8	43.4	51.5	18.2	46.7
- no bias	67.9	55.6	57.5	28.6	43.9	45.5	50.0	17.6	45.8
- no domain embedding init, no bias	70.0	55.4	56.1	25.9	41.1	44.9	44.9	16.6	44.4
- no training to coordinate	64.4	51.5	55.5	24.7	43.1	41.2	19.3	10.3	38.8
FLEXOLMO	71.0	59.8	59.9	30.8	45.8	47.1	50.7	17.3	47.8
<i>Reference model (upperbound)</i>									
Unrestricted MoE ($1 \times$ FLOPs, $0.5 \times$ Data)	68.0	53.8	57.8	28.9	41.5	48.6	49.4	22.2	46.3
Unrestricted MoE ($2 \times$ FLOPs, $1 \times$ Data)	73.3	60.2	63.1	32.5	48.1	54.4	53.4	27.0	51.5

Table 2: Evaluation of FLEXOLMO trained on *eight* sets (public mix and seven simulated closed sets) on 31 tasks across 10 categories, tested with 1,000 samples per subtask. “no RT” indicates no optional router training on proxy data (§3.3.3).

	MC9	GEN5	MMLU	MMLU Pro	AGIEval	BBH	Math2	NewsG	PoemG	SciRIFF5	Code4	Avg.
Prev. Public model	68.7	58.8	55.9	26.2	39.9	35.7	8.2	76.0	47.8	48.1	1.1	42.4
<i>Individual experts</i>												
Math	62.5	44.3	50.6	24.1	42.0	45.6	53.1	42.6	28.0	50.7	15.8	41.8
Code	40.5	39.4	29.5	14.5	27.4	38.1	6.0	45.1	28.2	48.0	21.0	30.7
Educational Text	64.3	52.1	56.5	27.0	39.7	40.3	13.6	57.6	51.8	51.7	3.0	41.6
News	46.5	48.6	36.4	15.2	25.7	30.9	2.5	77.7	26.9	47.0	0.0	32.5
Creative Writing	42.7	43.9	31.5	11.6	23.3	27.6	1.7	56.9	67.5	42.4	0.0	31.7
Academic	41.0	45.2	33.8	14.8	24.1	32.4	6.5	51.8	23.0	52.0	0.0	29.5
Reddit	64.7	36.5	56.1	25.5	35.5	19.7	2.5	54.1	8.6	32.7	1.7	30.7
<i>Combined model</i>												
BTM (top-2)	68.7	57.7	59.4	28.3	43.2	44.3	23.1	73.6	54.4	46.3	24.0	47.6
FLEXOLMO (no RT)	69.2	53.2	58.8	34.0	43.4	42.1	52.1	78.2	60.1	54.4	18.6	51.3
FLEXOLMO	70.8	59.8	60.4	30.9	45.1	46.4	48.5	80.7	62.2	54.3	17.2	52.4

individual expert performance, e.g., $35.6 \rightarrow 47.1$ on BBH, $8.1 \rightarrow 50.7$ on math, and $1.0 \rightarrow 17.3$ on coding. Notably, FLEXOLMO even matches or exceeds the performance of specialized experts on their respective tasks (e.g., on BBH and Math2).

FLEXOLMO achieves more effective merging than baselines We also compare FLEXOLMO with baseline merging methods (§4.3). All baselines outperform the model trained on Public Mix only. However, their performance is inconsistent: model soup and BTX are generally weak,³ while prompt-based routing is highly unstable: it performs well when the classifier selects the correct expert, but degrades sharply when it does not. Among the baselines, BTM yields the best performance. Nonetheless, FLEXOLMO outperforms all prior model merging methods, beating the best baseline BTM by 10.1% relative on average. We attribute this to the MoE-based design of our model, which selectively activates different experts per layer, effectively combining the complementary strengths of each specialized model (see further analysis in §5.2).

³This is likely because training on disjoint datasets causes experts to diverge from each other and from the seed model, making model soup limited, and training BTX on the public data only is not optimal.

Table 3: **Impact of embedding initialization methods on model performance.** The GRIT embedder [67] consistently outperforms public model embeddings across most benchmarks.

	Embed Init.	MC9	GEN5	MMLU	MMLU Pro	AGI Eval	BBH	Math2	Code4	Avg.
FLEXOLMO	Public	70.5	55.5	58.3	26.5	40.2	40.1	48.4	8.7	43.5
FLEXOLMO	GRIT	71.1	58.6	58.1	28.4	44.8	43.4	51.5	18.2	46.7




Figure 2: **Routing pattern analysis.** We visualize how text from different domains activate experts (four experts activated). The horizontal gray lines indicate uniform routing.

Comparison to the unrestricted MoE Compared to the unrestricted MoE trained without considering data restrictions (§4.3), FLEXOLMO outperforms the FLOP-controlled setting ($1 \times$ FLOPs, $0.5 \times$ Data). It slightly underperforms the data-controlled model ($2 \times$ FLOPs, $1 \times$ Data). This indicates that FLEXOLMO enables training without direct access to the data (requiring only model sharing) and flexible opt-in and opt-out functions *while* retaining strong performance.

Ablations on FLEXOLMO We further evaluate FLEXOLMO by removing different components introduced in §3.3: learning to coordinate, router initialization, and the bias term. Our results show that each component plays an important role, with the removal of any one leading to performance drop. In particular, we observe that randomly initializing router embeddings leads to the final learned router embeddings being very similar to each other, making the later merging of multiple experts harder. Furthermore, we confirm that FLEXOLMO benefits from additional router training (§3.3.3) and using external embedders for router initialization, compared to using the public model’s hidden states as the initialization (Table 3).

Final FLEXOLMO in the full setup Finally, we evaluate FLEXOLMO in the full eight-expert setup and compare it against the public-only model, individual experts trained on closed datasets, and BTM (top-2), the strongest baseline from Table 1. This was done by simply *adding four additional experts*, benefiting from FLEXOLMO’s flexibility in easily adding new datasets. Consistent with earlier findings, FLEXOLMO outperforms all individual models (the public baseline and individual experts), demonstrating the synergistic effect of combining independently trained modules. Compared to the strongest baseline BTM, it achieves a 10% relative improvement on average (Table 2). FLEXOLMO excels on benchmarks where specialized experts perform well (BBH, Math2, NewsG, PoemG, SciRIF5, Code4), matching or surpassing the experts, and also shows strong results on tasks where no single dataset suffices (e.g., MC9, Gen5, MMLU, MMLU Pro, AGI Eval).

5.2 Model Behavior Analysis

Routing patterns Figure 2 visualizes the router’s token distribution across experts for various domain inputs. The router tends to activate the corresponding domain expert (e.g., math inputs activate the math expert), demonstrating its ability to identify the most relevant module. We also observe frequent activation of the public expert, likely due to our coordinated training strategy, where each expert is designed to complement the public expert. Also, different combinations of experts are activated at different layers. This highlights the model’s layer-specific specialization and its greater expressivity than approaches that route inputs to a single expert (e.g., prompt-based routing).




Figure 3: Effect of active expert count




Figure 4: Opting out of news data.

Number of active experts We further analyze how the number of active experts affects downstream task performance. As shown in Figure 3, performance consistently improves as more experts are activated, up to four experts, after which it plateaus. This suggests that the final model can operate efficiently as a sparse model by activating only four experts per input during inference.

Data opt-out FLEXOLMO offers a straightforward mechanism for opting out of specific datasets by removing the corresponding expert module at inference time. In Figure 4, we evaluate the model’s performance after excluding the news expert. As expected, performance drops on in-domain tasks such as news generation. However, on unrelated tasks, performance remains largely unaffected.

5.3 Data Extraction Analysis

FLEXOLMO enables data owners to contribute to the model without sharing their data by instead sharing model weights trained locally on their data. A natural and important concern is whether their data can be extracted from these shared weights [108, 109, 110]. This risk is particularly relevant when training data includes private or confidential information.

We empirically assess this risk by implementing training data extraction following prior work [108]. Specifically, we sample 10,000 documents from the math data.⁴ From each document, we extract a 32-token prefix and use it to generate 256-token continuation using top- k sampling ($k = 50$), top- p sampling ($p = 0.95$), and a temperature of 1.0. We sample 10 times per prefix, and if any of the generated outputs achieves a normalized Levenshtein similarity of 0.9 or higher with the original document, we consider that document to be extracted. To validate our implementation, we apply it to a model overfitted on the dataset (trained for 100 epochs), and observe a 60% extraction rate.

Our results are as follows: (1) A public model that has not seen any math data yields an extraction rate of 0.1%. (2) A dense model trained on the math dataset (i.e., math expert) yields 1.6%. (3) FLEXOLMO with the math expert included yields 0.7%.

These results lead to the following conclusions. First, in practice, it is difficult to extract a substantial portion of the training data, which is in line with previous findings [108]. However, if a model includes any weights trained on the data, nonzero (though small) fraction of the data may be extractable. If data owners are comfortable with this minimal leakage, as long as a meaningful fraction of the data remains not extractable, we believe that FLEXOLMO, in its current form, is a viable solution. If the owners’ data includes any private or sensitive information, we recommend training experts using differentially private (DP) learning methods before contributing them to the model, which provides formal privacy guarantees. Applying DP is orthogonal to our architecture, and different data owners can make independent decisions, providing flexibility without compromising the overall design.

6 Conclusion

We introduce FLEXOLMO, a new class of LMs that solves real-world data constraint challenges with (1) *modular, distributed training*, where different model parameters are independently trained on disjoint and locally maintained datasets, and (2) *data-flexible inference*, where data can be selectively included at inference-time, with guarantees. We show that FLEXOLMO significantly outperforms competitive baselines, while providing the benefits of distributed training and flexible inference.

⁴We chose the math data because it is the smallest among our simulated closed sets and the math expert is trained for three epochs (instead of one), making it more susceptible to extraction. Therefore, the extraction rates with the math data likely represent an upper bound.

References

- [1] Pratyush Maini, Zhili Feng, Avi Schwarzschild, Zachary Chase Lipton, and J Zico Kolter. TOFU: A task of fictitious unlearning for LLMs. In *First Conference on Language Modeling*, 2024.
- [2] [REDACTED], Jaechan Lee, Yangsibo Huang, Sadhika Malladi, Jieyu Zhao, Ari Holtzman, [REDACTED], [REDACTED], A [REDACTED], and Chiyuan Zhang. Muse: Machine unlearning six-way evaluation for language models. *arXiv preprint [REDACTED]*, 2024.
- [3] A Feder Cooper, Christopher A Choquette-Choo, Miranda Bogen, Matthew Jagielski, Katja Filippova, Ken Ziyu [REDACTED], Alexandra Chouldechova, Jamie Hayes, Yangsibo Huang, Niloofar Mireshghallah, et al. Machine unlearning doesn’t do what you think: Lessons for generative ai policy, research, and practice. *arXiv preprint [REDACTED]*, 2024.
- [4] Jakub Konecný, H Brendan McMahan, X Yu Felix, Peter Richtárik, Ananda Theertha Suresh, and Dave Bacon. Federated learning: Strategies for improving communication efficiency. *CoRR*, 2016.
- [5] Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera y Arcas. Communication-Efficient Learning of Deep Networks from Decentralized Data. In Aarti Singh and Jerry Zhu, editors, *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, volume 54 of *Proceedings of Machine Learning Research*, pages 1273–1282. PMLR, 20–22 Apr 2017.
- [6] Kang Wei, Jun Li, [REDACTED]g Ding, Chuan Ma, Howard Hua Yang, Farokhi Farhad, [REDACTED] Jin, Tony Q. S. Quek, and H. Vincent Poor. Federated learning with differential privacy: Algorithms and performance analysis. *IEEE Transactions on Information Forensics and Security*, 15:3454–3469, 2019.
- [7] Xiaojin Zhang, Yan Kang, Kai Chen, Lixin Fan, and Qiang Yang. Trading off privacy, utility, and efficiency in federated learning. *ACM Transactions on Intelligent Systems and Technology*, 14:1 – 32, 2022.
- [8] Noam Shazeer, Az [REDACTED]a Mirhoseini, Krzysztof Maziarsz, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. *arXiv preprint [REDACTED]*, 2017.
- [9] [REDACTED], [REDACTED], [REDACTED], [REDACTED], [REDACTED], [REDACTED], [REDACTED], Pete [REDACTED], Oyvind Tafjord, Nathan Lambert, et al. Olmoe: Open mixture-of-experts language models. *arXiv preprint [REDACTED]*, 2024.
- [10] Prateek Yadav, Colin Raffel, Mohammed Muqeeth, [REDACTED]s Caccia, Haokun [REDACTED], Tianlong Chen, Mohit Bansal, Leshem Choshen, and Alessandro Sordoni. A survey on model moering: Recycling and routing among specialized experts for collaborative learning. *CoRR*, 2024.
- [11] [REDACTED], Suchin Gururangan, Tim Dettmers, [REDACTED], Tim Althoff, [REDACTED], and [REDACTED] Branch-train-merge: Embarrassingly parallel training of expert language models, 2022.
- [12] Mitchell Wortsman, Gabriel Ilharco, Samir Ya Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes, Ari S Morcos, Hongseok Namkoong, [REDACTED], Yair Carmon, Simon Kornblith,

- et al. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In *International conference on machine learning*, pages 23965–23998. PMLR, 2022.
- [13] [REDACTED], Hongjin Su, Liang Wang, Nan Yang, Furu Wei, Tao Yu, Amanpreet Singh, and Douwe Kiela. Generative representational instruction tuning, 2024.
 - [14] BIS Innovation Hub Nordic Centre. Project aurora: the power of data, technology and collaboration to combat money laundering across institutions and borders. Technical report. Occasional publication No. 66.
 - [15] Chen Zhang, Yu Xie, Hang Bai, Bin Yu, Weihong Li, and Yuan Gao. A survey on federated learning. *Knowledge-Based Systems*, 216:106775, 2021.
 - [16] Patrik Hummel, Matthias Braun, Max Tretter, and Peter Dabrock. Data sovereignty: A review. *Big Data & Society*, 8(1):2053951720982012, 2021.
 - [17] Martin Grund, Stefania Leone, Herman van Hövell, Sven Wagner-Boysen, Sebastian Hillig, Hyukjin Kwon, David [REDACTED], Jakob Mund, Polo-Francois Poli, Lionel Montrieux, et al. Databricks lakeguard: Supporting fine-grained access control and multi-user capabilities for apache spark workloads. In *Companion of the 2025 International Conference on Management of Data*, pages 418–430, 2025.
 - [18] Albert Ziegler, Eirini Kalliamvakou, X [REDACTED] Li, Andrew Rice, Devon Rifkin, Shawn Simister, Ganesh Sittampalam, and Edward Aftandilian. Productivity assessment of neural code completion. In *Proceedings of the 6th ACM SIGPLAN International Symposium on Machine Programming*, pages 21–29, 2022.
 - [19] Peter Henderson, Xuechen Li, Dan Jurafsky, Tatsunori Ha [REDACTED] moto, Mark A Lemley, and Percy Liang. Foundation models and fair use. *Journal of Machine Learning Research*, 24(400):1–79, 2023.
 - [20] [REDACTED], Robert Mahari, Ariel Lee, Campbell Lund, Hamidah Oderinwale, William Brannon, Nayan Saxena, Naana Obeng-Marnu, Tobin South, Cole Hunter, et al. Consent in crisis: The rapid decline of the ai data commons. *Advances in Neural Information Processing Systems*, 37:108042–108087, 2024.
 - [21] Ziv Epstein, Aaron Hertzmann, Investigators of Human Creativity, Memo Akten, Hany Farid, Jessica Fjeld, Morgan R Frank, Matthew Groh, Laura Herman, Neil Leach, et al. Art and the science of generative ai. *Science*, 380(6650):1110–1111, 2023.
 - [22] Ri [REDACTED] Bommasani, [REDACTED] Klyman, [REDACTED], Sayash Kapoor, Nestor Maslej, Betty Xiong, Daniel Zhang, and Percy Liang. The foundation model transparency index. *arXiv preprint [REDACTED]*, 2023.
 - [23] Peter Henderson, Xuechen Li, Dan Jurafsky, Tatsunori Ha [REDACTED] moto, Mark A Lemley, and Percy Liang. Foundation models and copyright questions, 2023.
 - [24] [REDACTED], Robert Mahari, Anthony Chen, Naana Obeng-Marnu, Damien Sileo, William Brannon, [REDACTED], Nathan Khazam, Jad Kabbara, Kartik Perisetla, Xinyi Wu, Enrico [REDACTED] ppole, Kurt D. Bollacker, Tongshuang Wu, Luis Villa, Sandy Pentland, and Sara Hooker. A large-scale audit of dataset licensing and attribution in ai. *Nat. Mac. Intell.*, 6(8):975–987, 2024.
 - [25] Ri [REDACTED] Bommasani, Drew A Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, et al. On the opportunities and risks of foundation models. *arXiv preprint [REDACTED]*, 2021.
 - [26] [REDACTED], Robert Mahari, Anthony Chen, Naana Obeng-Marnu, Damien Sileo, William Brannon, [REDACTED], Nathan Khazam, Jad Kabbara, Kartik Perisetla, et al. The data provenance initiative: A large scale audit of dataset licensing & attribution in ai. *arXiv preprint [REDACTED]*, 2023.

- [27] Abeba Birhane, Sanghyun Han, Vishnu Boddeti, Sasha Luccioni, et al. Into the laion’s den: Investigating hate in multimodal datasets. *Advances in neural information processing systems*, 36:21268–21284, 2023.
- [28] Jesse Dodge, Maarten Sap, Ana Marasović, William Agnew, Gabriel Ilharco, [REDACTED], [REDACTED] Mitchell, and Matt Gardner. Documenting large webtext corpora: A case study on the colossal clean crawled corpus. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 1286–1305, 2021.
- [29] [REDACTED], Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, [REDACTED], Terra Blevins, Danqi Chen, and [REDACTED]. Detecting pretraining data from large language models. In *The Twelfth International Conference on Learning Representations*.
- [30] Peter Kairouz, H Brendan McMahan, Brendan Avent, Aurélien Bellet, Mehdi Bennis, Arjun Nitin Bhagoji, Kallista Bonawitz, Zachary Charles, Graham Cormode, Rachel Cummings, et al. Advances and open problems in federated learning. *Foundations and trends® in machine learning*, 14(1–2):1–210, 2021.
- [31] Tian Li, Anit Kumar Sahu, Manzil Zaheer, Maziar Sanjabi, Ameet Talwalkar, and Virginia [REDACTED]. Federated optimization in heterogeneous networks. *Proceedings of Machine learning and systems*, 2:429–450, 2020.
- [32] Xinghao Wu, Xuefeng [REDACTED], Jianwei Niu, Haolin Wang, Shaojie Tang, and Guogang Zhu. Fedlora: When personalized federated learning meets low-rank adaptation. 2024.
- [33] Haodong Zhao, Wei Du, Fangqi Li, Peixuan Li, and Gongshen [REDACTED]. Fedprompt: Communication-efficient and privacy-preserving prompt tuning in federated learning. In *ICASSP 2023-2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5. IEEE, 2023.
- [34] [REDACTED] Kuo, Arian Raje, Kousik Rajesh, and Virginia [REDACTED]. Federated lora with sparse communication. *arXiv preprint [REDACTED]*, 2024.
- [35] Pascal Paillier. Public-key cryptosystems based on composite degree residuosity classes. In *Proceedings of the 17th International Conference on Theory and Application of Cryptographic Techniques, EUROCRYPT’99*, page 223–238, Berlin, Heidelberg, 1999. Springer-Verlag.
- [36] Cynthia Dwork, Aaron Roth, et al. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science*, 9(3–4):211–407, 2014.
- [37] Hongyi Wang, Kartik Sreenivasan, Shashank Rajput, Harit Vishwakarma, Saurabh Agarwal, Jy-yong Sohn, Kangwook Lee, and Dimitris Papailiopoulos. Attack of the tails: Yes, you really can backdoor federated learning. In H. Larochelle, M. Ranzato, R. Hadsell, M.F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 16070–16084. Curran Associates, Inc., 2020.
- [38] Erfan Darzi, Florian Dubost, Nanna M Sijtsma, and Peter MA van Ooijen. Exploring adversarial attacks in federated learning for medical imaging. *IEEE Transactions on Industrial Informatics*, 2024.
- [39] John Nguyen, K[REDACTED]tiz M[REDACTED]k, Hongyuan Zhan, Ashkan Yousefpour, [REDACTED] Rabbat, Mani Malek, and Dzmitry Huba. Federated learning with buffered asynchronous aggregation. In *International conference on artificial intelligence and statistics*, pages 3581–3607. PMLR, 2022.
- [40] Colin Raffel. Building machine learning models like open source software. *Communications of the ACM*, 66(2):38–40, 2023.
- [41] Shangbin Feng, Zifeng Wang, Palash Goyal, Yike Wang, [REDACTED], Huang Xia, Hamid Palangi, [REDACTED], Yulia Tsvetkov, Chen-Yu Lee, and Tomas Pfister. Heterogeneous swarms: Jointly optimizing model roles and weights for multi-llm systems, 2025.

- [42] Alexandra Chronopoulou, Matthew E Peters, Alexander Fraser, and Jesse Dodge. Adaptersoup: Weight averaging to improve generalization of pretrained language models. *arXiv preprint* [arXiv:2305.10345](#), 2023.
- [43] Jyothish Pari, Samy Jelassi, and Pulkit Agrawal. Collective model intelligence requires compatible specialization. *ArXiv*, abs/2411.02207, 2024.
- [44] Shangbin Feng, Zifeng Wang, Yike Wang, Sayna Ebrahimi, Hamid Palangi, Lesly Miculicich, Achin Kulshrestha, Nathaniel Rauschmayr, Yejin Choi, Yulia Tsvetkov, et al. Model swarms: Collaborative search to adapt llm experts via swarm intelligence. *arXiv e-prints*, pages arXiv–2410, 2024.
- [45] Prateek Yadav, Derek Tam, Leshem Choshen, Colin A Raffel, and Mohit Bansal. Ties-merging: Resolving interference when merging models. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 7093–7115. Curran Associates, Inc., 2023.
- [46] Michael S Matena and Colin A Raffel. Merging models with fisher-weighted averaging. In S. Koyejo, S. Mohamed, A. Agarwal, D. Belgrave, K. Cho, and A. Oh, editors, *Advances in Neural Information Processing Systems*, volume 35, pages 17703–17716. Curran Associates, Inc., 2022.
- [47] Gabriel Ilharco, Marco Tulio Ribeiro, Mitchell Wortsman, Ludwig Schmidt, Hanja Jia, and Ludwig Schmidt. Editing models with task arithmetic. In *The Eleventh International Conference on Learning Representations*, 2023.
- [48] Takuya Akiba, Makoto Tanaka, Yujin Tang, Qi Sun, and David Ha. Evolutionary optimization of model merging recipes. *arXiv preprint* [arXiv:2405.14272](#), 2024.
- [49] Enneng Yang, Zhenyi Wang, Li Shen, Xingwei Guo, Guibing Guo, Xingwei Wang, and Dacheng Tao. Adamerging: Adaptive model merging for multi-task learning. In *The Twelfth International Conference on Learning Representations*, 2024.
- [50] [Jesse Dodge](#), [Pradeep Dasigi](#), [Jesse Dodge](#), [Pradeep Dasigi](#), Jesse Dodge, and Pradeep Dasigi. Merge to learn: Efficiently adding skills to language models with model merging, 2024.
- [51] Shangbin Feng, Wenxuan Ding, Zifeng Wang, Yike Wang, Zejiang Shen, Xiaochuang Han, Hunter Lang, Chen-Yu Lee, Tomas Pfister, Yejin Choi, and Yulia Tsvetkov. When one llm drools, multi-llm collaboration rules, 2025.
- [52] [Suchin Gururangan](#), [Tim Dettmers](#), [Tim Althoff](#), [A](#), and [Suchin Gururangan](#). Branch-train-merge: Embarrassingly parallel training of expert language models. *arXiv preprint* [arXiv:2205.12244](#), 2022.
- [53] Suchin Gururangan, [Suchin Gururangan](#), [Tim Althoff](#), [A](#), and [Suchin Gururangan](#). Scaling expert language models with unsupervised domain discovery. *arXiv preprint* [arXiv:2305.12244](#), 2023.
- [54] Suchin Gururangan, [Ari Holtzman](#), [A](#), and [Suchin Gururangan](#). Demix layers: Disentangling domains for modular language modeling. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 5557–5576, 2022.
- [55] Sainbayar Sukhbaatar, Olga Golovneva, Vasu Sharma, Hu Xu, Xi Victoria Lin, Baptiste Rozière, [Kahn](#), Daniel Li, [Jason Weston](#), et al. Branch-train-mix: Mixing expert llms into a mixture-of-experts llm. *arXiv preprint* [arXiv:2405.12244](#), 2024.
- [56] Peter Schafhalter, Shun Liao, Yanqi Zhou, Chih-Kuan Yeh, Arun Kandoor, and James Laudon. Scalable multi-domain adaptation of language models using modular experts. *arXiv preprint* [arXiv:2405.12244](#), 2024.

- [57] Qizhen Zhang, Prajjwal Bhargava, Chloe Bi, Chris X Cai, Jakob Foerster, Jeremy Fu, Punit Singh Koura, Ruan Silva, Sheng Shen, Emily Dinan, et al. Bts: Harmonizing specialized experts into a generalist llm. *arXiv preprint* [REDACTED], 2025.
- [58] Qizhen Zhang, Nikolas Gritsch, Dwaraknath Gnaneshwar, Simon Guo, David Cairuz, Bharat Venkitesh, Jakob Foerster, Phil Blunsom, Sebastian Ruder, A. Ustun, and Acyr F. Locatelli. Bam! just like that: Simple and efficient parameter upcycling for mixture of experts. *ArXiv*, abs/2408.08274, 2024.
- [59] Yikang Shen, Zheyu Zhang, Tianyou Cao, Shawn Tan, Zhenfang Chen, and Chuang Gan. Moduleformer: Modularity emerges from mixture-of-experts. *arXiv preprint* [REDACTED], 2023.
- [60] Ted Zadouri, Ahmet Üstün, Arash Ahmadian, Beyza Ermiş, Acyr Locatelli, and Sara Hooker. Putting mixture of experts to the limit: Extremely parameter efficient moe for instruction tuning. *arXiv preprint* [REDACTED], 2023.
- [61] Yun Zhu, N[REDACTED] Wichers, Chu-Cheng Lin, Xinyi Wang, Tianlong Chen, Lei Shu, Han Lu, Canoe [REDACTED], Liangchen Luo, Jindong Chen, and Lei Meng. Sira: Sparse mixture of low rank adaptation, 2023.
- [62] [REDACTED]han Dou, Enyu Zhou, Yan [REDACTED], Songyang Gao, Wei Shen, Limao Xiong, Yuhao Zhou, Xiao Wang, Zhiheng Xi, Xiaoran Fan, [REDACTED]liang Pu, Jiang Zhu, Rui Zheng, Tao Gui, Qi Zhang, and Xuanjing Huang. LoRAMoE: Alleviating world knowledge forgetting in large language models via MoE-style plugin. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1932–1945, Bangkok, Thailand, August 2024. Association for Computational Linguistics.
- [63] Jonas Pfeiffer, Aishwarya Kamath, Andreas Rücklé, Kyunghyun Cho, and Iryna Gurevych. Adapterfusion: Non-destructive task composition for transfer learning. *arXiv preprint* [REDACTED], 2020.
- [64] William Fleshman, Aleem Khan, Marc Marone, and Benja[REDACTED] Van Durme. Adapterswap: Continuous training of llms with data removal and access-control guarantees. *CoRR*, 2024.
- [65] [REDACTED] Kuo, Amrith Setlur, Kartik Srinivas, Aditi Raghunathan, and Virginia [REDACTED]. Exact unlearning of finetuning data via model merging at scale. *arXiv preprint* [REDACTED], 2025.
- [66] Damai Dai, Chengqi Deng, Chenggang Zhao, RX Xu, Huazuo Gao, Deli Chen, Jia [REDACTED] Li, Wangding Zeng, Xingkai Yu, Yu Wu, et al. Deepseekmoe: Towards ultimate expert specialization in mixture-of-experts language models. *arXiv preprint* [REDACTED], 2024.
- [67] [REDACTED], Hongjin SU, Liang Wang, Nan Yang, Furu Wei, Tao Yu, Amanpreet Singh, and Douwe Kiela. Generative representational instruction tuning. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [68] Nikolas Gritsch, Qizhen Zhang, Acyr Locatelli, Sara Hooker, and Ahmet Üstün. Nexus: Specialization meets adaptability for efficiently training mixture of experts. *arXiv preprint* [REDACTED], 2024.
- [69] Ziyu Zhao, Leilei Gan, Guoyin Wang, Wangchunshu Zhou, Hongxia Yang, Kun Kuang, and Fei Wu. Loraretriever: Input-aware lora retrieval and composition for mixed tasks in the wild. *arXiv preprint* [REDACTED], 2024.
- [70] Joel Jang, Seungone Kim, Seonghyeon Ye, Doyoung Kim, Lajanugen Logeswaran, Moontae Lee, Kyungjae Lee, and [REDACTED]joon Seo. Exploring the benefits of training expert language models over instruction tuning. In *International Conference on Machine Learning*, pages 14702–14729. PMLR, 2023.
- [71] Feng Cheng, Ziyang Wang, Yi-Lin Sung, Yan-Bo Lin, Mohit Bansal, and Gedas Bertasius. Dam: Dynamic adapter merging for continual video qa learning. In *2025 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 6805–6817. IEEE, 2025.

- [72] Hyunji Lee, [REDACTED], Arman Cohan, [REDACTED] Joon Seo, and [REDACTED]. Routerretriever: Routing over a mixture of expert embedding models. *arXiv preprint* [REDACTED], 2024.
- [73] Joshua Belofsky. Token-level adaptation of lora adapters for downstream task generalization. In *Proceedings of the 2023 6th Artificial Intelligence and Cloud Computing Conference*, pages 168–172, 2023.
- [74] Jeffrey Li, Alex Fang, Georgios Smyrnis, Maor Ivgi, Matt Jordan, Samir Gadre, Hritik Bansal, Etash Guha, Sedrick Keh, Kushal Arora, Saurabh Garg, Rui Xin, [REDACTED], Reinhard Heckel, Jean Mercat, Mayee Chen, Suchin Gururangan, Mitchell Wortsman, Alon Albalak, Yonatan Bitton, Marianna Nezhurina, Amro Abbas, Cheng-Yu Hsieh, Dhruva Ghosh, Josh Gardner, Maciej Kilian, Hanlin Zhang, Rulin Shao, Sarah Pratt, Sunny Sanyal, Gabriel Ilharco, Giannis Daras, Kalyani Marathe, Aaron Gokaslan, Jieyu Zhang, Khyathi Chandu, Thao Nguyen, Igor Vasiljevic, Sham Kakade, Shuran Song, Sujay Sanghavi, Fartash Faghri, Sewoong Oh, [REDACTED], [REDACTED], Alaaeldin El-Nouby, Hadi Pouransari, Alexander Toshev, Stephanie Wang, [REDACTED], [REDACTED], [REDACTED], Jenia Jitsev, Thomas Kollar, Alexandros G. Dimakis, Yair Carmon, Achal Dave, Ludwig Schmidt, and Vaishaal Shankar. Datacomp-lm: In search of the next generation of training sets for language models, 2025.
- [75] Alexander Wettig, [REDACTED], [REDACTED], [REDACTED], Danqi Chen, and [REDACTED]. Organize the web: Constructing domains enhances pre-training data curation, 2025.
- [76] Raymond Li, Loubna Ben Allal, Yangtian Zi, [REDACTED], Denis Kocetkov, Chenghao Mou, Marc Marone, Christopher Akiki, Jia Li, Jenny Chim, Qian [REDACTED], Evgenii Zheltonozhskii, Terry Yue Zhuo, Thomas Wang, Olivier Dehaene, Ming Davaadorj, Joel Lamy-Poirier, João Monteiro, Oleh Shliazhko, Nicolas Gontier, Nicholas Meade, Armel Zebaze, [REDACTED]-Ho Yee, Logesh Kumar Umapathi, Jian Zhu, Benja [REDACTED] Lipkin, Muhtasham Oblokulov, Zhiruo Wang, Rudra Murthy, Jason Stillerman, Siva Sankalp Patel, Dmitry Abulkhanov, Marco Zocca, Manan Dey, Zhihan Zhang, Nour Fahmy, Urva [REDACTED] Bhattacharyya, Wenhao Yu, Swayam Singh, Sasha Luccioni, Paulo Villegas, Maxim Kunakov, Fedor Zhdanov, Manuel Romero, Tony Lee, Nadav Timor, Jennifer Ding, Claire Schlesinger, Hailey Schoelkopf, Jan Ebert, Tri Dao, Mayank Mishra, Alex Gu, Jennifer Robinson, Carolyn Jane Anderson, Brendan Dolan-Gavitt, Danish Contractor, Siva Reddy, Daniel Fried, Dzmitry Bahdanau, Yacine Jernite, Carlos Muñoz Ferrandis, Sean Hughes, Thomas Wolf, Arjun Guha, Leandro von Werra, and Harm de Vries. Starcoder: may the source be with you!, 2023.
- [77] Denis Kocetkov, Raymond Li, Loubna Ben Allal, Jia Li, Chenghao Mou, Carlos Muñoz Ferrandis, Yacine Jernite, [REDACTED] Mitchell, Sean Hughes, Thomas Wolf, Dzmitry Bahdanau, Leandro von Werra, and Harm de Vries. The stack: 3 tb of permissively licensed source code, 2022.
- [78] Team OLMo, Pete [REDACTED], [REDACTED], [REDACTED], [REDACTED], Shane Arora, [REDACTED], [REDACTED], Yuling Gu, Shengyi Huang, Matt Jordan, Nathan Lambert, [REDACTED], Oyvind Tafjord, Taira Anderson, David Atkinson, Faeze Brahman, Christopher Clark, Pradeep Dasigi, Nouha Dziri, Michal Guerquin, Hamish Ivison, [REDACTED], Jiacheng [REDACTED], Saumya M [REDACTED], William Merrill, Lester James V. Miranda, [REDACTED], Tyler Murray, Crystal Nam, Valentina Pyatkin, Aman Rangapur, Michael Schmitz, Sam Skjonsberg, David Wadden, Christopher Wilhelm, Michael Wilson, [REDACTED], [REDACTED], [REDACTED], and [REDACTED]. 2 olmo 2 furious, 2025.
- [79] [REDACTED], Jon Borchardt, Jason Dunkelberger, Regan Huff, Daniel Lin, Aman Rangapur, Christopher Wilhelm, [REDACTED], and [REDACTED]. olmocr: Unlocking trillions of tokens in pdfs with vision language models, 2025.
- [80] [REDACTED] and [REDACTED]. peS2o (Pretraining Efficiently on S2ORC) Dataset. Technical report, Allen Institute for AI, 2023. ODC-By, [REDACTED]
- [81] [REDACTED], Lucy Lu Wang, Mark Neumann, Rodney Kinney, and Daniel Weld. S2ORC: The semantic scholar open research corpus. In Dan Jurafsky, Joyce Chai, Nat [REDACTED] Schluter, and Joel Tetreault, editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4969–4983, Online, July 2020. Association for Computational Linguistics.

- [82] Loubna Ben Allal, Anton Lozhkov, Elie Bakouch, Gabriel Martín Blázquez, Guilherme Penedo, [REDACTED] Tunstall, Andrés Marafioti, Hynek Kydlíček, Agustín Piqueres Lajarín, Vaibhav Srivastav, Joshua Lochner, Caleb Fahlgren, Xuan-Son Nguyen, Clémentine Fourier, Ben Burtenshaw, Hugo Larcher, Haojun Zhao, Cyril Zakka, Mathieu Morlon, Colin Raffel, Leandro von Werra, and Thomas Wolf. Smollm2: When smol goes big – data-centric training of a small language model, 2025.
- [83] [REDACTED], Rodney Kinney, [REDACTED], [REDACTED], David Atkinson, Russell Authur, Ben Bogin, Khyathi Chandu, Jennifer Dumas, Yanai Elazar, Valentin Hofmann, Ananya Harsh Jha, Sachin Kumar, Li Lucy, Xinxu Lyu, Nathan Lambert, Ian Magnusson, [REDACTED], [REDACTED], Aakanksha Naik, Crystal Nam, Matthew E. Peters, Abhilasha Ravichander, [REDACTED] Richardson, Zejiang Shen, Emma Strubell, Nishant Subramani, Oyvind Tafjord, Pete [REDACTED], [REDACTED], [REDACTED], [REDACTED], Iz Beltagy, [REDACTED], [REDACTED], Jesse Dodge, and [REDACTED]. Dolma: an open corpus of three trillion tokens for language model pretraining research, 2024.
- [84] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, A[REDACTED]sh Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge. [REDACTED], 2018.
- [85] Christopher Clark, Kenton Lee, [REDACTED]-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. Boolq: Exploring the surprising difficulty of natural yes/no questions, 2019.
- [86] Amrita Saha, Vardaan Pahuja, Mitesh M. Khapra, Karthik Sankaranarayanan, and Sarath Chandar. Complex sequential question answering: Towards learning to converse over linked question answer pairs with a knowledge graph, 2018.
- [87] Rowan Zellers, Ari Holtzman, Yonatan Bisk, [REDACTED], and Yejin Choi. Hellaswag: Can a machine really finish your sentence?, 2019.
- [88] Todor Mihaylov, Peter Clark, Tushar Khot, and A[REDACTED]sh Sabharwal. Can a suit of armor conduct electricity? a new dataset for open book question answering. In *EMNLP*, 2018.
- [89] Yonatan Bisk, Rowan Zellers, Ronan Le Bras, Jianfeng Gao, and Yejin Choi. Piqa: Reasoning about physical commonsense in natural language, 2019.
- [90] Maarten Sap, Hannah Rashkin, Derek Chen, Ronan LeBras, and Yejin Choi. Soci[REDACTED]qa: Commonsense reasoning about social interactions, 2019.
- [91] Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. Winogrande: An adversarial winograd schema challenge at scale. *arXiv preprint* [REDACTED], 2019.
- [92] Siva Reddy, Danqi Chen, and Christopher D. Manning. Coqa: A conversational question answering challenge, 2019.
- [93] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. Squad: 100,000+ questions for machine comprehension of text, 2016.
- [94] Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, [REDACTED] Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, [REDACTED]-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. Natural questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:452–466, 2019.
- [95] Mandar Jo[REDACTED], Eunsol Choi, Daniel S. Weld, and [REDACTED] Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension, 2017.
- [96] Dheeru Dua, Yizhong Wang, Pradeep Dasigi, Gabriel Stanovsky, Sameer Singh, and Matt Gardner. Drop: A reading comprehension benchmark requiring discrete reasoning over paragraphs, 2019.

- [97] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and [REDACTED] Steinhardt. Measuring massive multitask language understanding. *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.
- [98] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, [REDACTED]guang Guo, Wei [REDACTED]g Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, and Wenhui Chen. Mmlu-pro: A more robust and challenging multi-task language understanding benchmark, 2024.
- [99] Wanjun Zhong, Ruixiang Cui, Yiduo Guo, Yaobo Liang, Shuai Lu, Yanlin Wang, A [REDACTED] Saied, Weizhu Chen, and Nan Duan. Agieval: A human-centric benchmark for evaluating foundation models, 2023.
- [100] Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc V. Le, Ed H. Chi, Denny Zhou, and Jason Wei. Challenging big-bench tasks and whether chain-of-thought can solve them, 2022.
- [101] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, [REDACTED] Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems, 2021.
- [102] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and [REDACTED] Steinhardt. Measuring mathematical problem solving with the math dataset, 2021.
- [103] [REDACTED] Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. Program synthesis with large language models, 2021.
- [104] Jiawei [REDACTED], Chunqiu Steven Xia, Yuyao Wang, and LING [REDACTED]G ZHANG. Is your code generated by chatGPT really correct? rigorous evaluation of large language models for code generation. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.
- [105] Mark Chen, Jerry Tworek, Heewoo Jun, Qi [REDACTED]g Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, [REDACTED] Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code, 2021.
- [106] David Wadden, Kejian [REDACTED], [REDACTED], Aakanksha Naik, Shruti Singh, Nitzan Barzilay, [REDACTED], Tom Hope, [REDACTED], Shannon Zejiang Shen, Doug Downey, [REDACTED], and Arman Cohan. Sciriff: A resource to enhance language model instruction-following over scientific literature, 2024.
- [107] Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, Anirudh Goyal, Anthony S. Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aur’elien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Rozière, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Cantón Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab A. AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael [REDACTED], Filip Radenovic, Frank Zhang, Gabriele Synnaeve, Gabrielle Lee, Georgia [REDACTED] Anderson, Graeme Nail, Grégoire Mialon, Guanglong [REDACTED], Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan

- Zarov, Imanol Arrieta Ibarra, Isabel M. Kloumann, Ishan Misra, Ivan Evtimov, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen [REDACTED], Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua Saxe, Ju-Qing Jia, Kalyan Vasuden Alwala, K. Upasani, Kate Plawiak, Keqian Li, Ken-591 neth Heafield, [REDACTED] R. Stone, Kh[REDACTED]d El-Arini, Krithika Iyer, K[REDACTED]tiz M[REDACTED]k, Kuen ley Chiu, Kunal Bhalla, Lauren Rantala-Yearly, Laurens van der Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo, Lukas Blecher, Lukas Landzaat, [REDACTED] de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat Singh, Manohar Paluri, Marcin Kardas, Mathew Oldham, Mathieu Rita, Maya Pavlova, Melissa Hall Melanie Kambadur, [REDACTED], (...), Yu Wang, Yuchen Hao, Yundi Qian, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu Yang, and Zhiwei Zhao. The llama 3 herd of models. *ArXiv*, abs/2407.21783, 2024.
- [108] Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pages 2633–2650, 2021.
- [109] Nicolas Carlini, Jamie Hayes, Milad Nasr, Matthew Jagielski, Vikash Sehwal, Florian Tramer, Borja Balle, Daphne Ippolito, and Eric Wallace. Extracting training data from diffusion models. In *32nd USENIX Security Symposium (USENIX Security 23)*, pages 5253–5270, 2023.
- [110] A Feder Cooper, Aaron Gokaslan, Amy B Cyphert, Christopher De Sa, Mark A Lemley, Daniel E Ho, and Percy Liang. Extracting memorized pieces of (copyrighted) books from open-weight language models. *arXiv preprint [REDACTED]*, 2025.
- [111] [REDACTED] Merrick, Danmei Xu, Gaurav Nuti, and Daniel Campos. Arctic-embed: Scalable, efficient, and accurate text embedding models. *arXiv preprint [REDACTED]*, 2024.
- [112] Jason Baumgartner, Savvas Zannettou, Brian Keegan, Megan Squire, and Jeremy Blackburn. The push [REDACTED]ft reddit dataset. In *Proceedings of the international AAAI conference on web and social media*, volume 14, pages 830–839, 2020.
- [113] [REDACTED], Iz Beltagy, [REDACTED], [REDACTED], Rodney Kinney, Oyvind Tafjord, Ananya Jha, Hamish Ivison, Ian Magnusson, Yizhong Wang, Shane Arora, David Atkinson, Russell Authur, Khyathi Chandu, Arman Cohan, Jennifer Dumas, Yanai Elazar, Yuling Gu, Jack Hessel, Tushar Khot, William Merrill, [REDACTED], [REDACTED], Aakanksha Naik, Crystal Nam, Matthew Peters, Valentina Pyatkin, Abhilasha Ravichander, [REDACTED], Saurabh Shah, William [REDACTED], Emma Strubell, Nishant Subramani, Mitchell Wortsman, Pradeep Dasigi, Nathan Lambert, [REDACTED] Richardson, [REDACTED] Jesse Dodge, [REDACTED], [REDACTED], [REDACTED], and [REDACTED]. OLMo: Accelerating the science of language models. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15789–15809, Bangkok, Thailand, August 2024. Association for Computational Linguistics.
- [114] Cody Blakeney, Mansheej Paul, Brett W. Larsen, Sean Owen, and Jonathan Frankle. Does your data spark joy? performance gains from domain upsampling at the end of training, 2024.
- [115] Tomáš Kočiský, Jonathan Schwarz, Phil Blunsom, Chris Dyer, Karl Moritz Hermann, Gábor Melis, and Edward Grefenstette. The narrativeqa reading comprehension challenge, 2017.
- [116] Yuling Gu, Oyvind Tafjord, Bailey Kuehl, Dany Haddad, Jesse Dodge, and [REDACTED] Ha-jil [REDACTED]rzi. Olmes: A standard for language model evaluations, 2025.
- [117] Aitor Lewkowycz, Anders Andreassen, David Dohan, Ethan Dyer, Henryk Michalewski, Vinay Ramasesh, Ambrose Slone, Cem Anil, Imanol Schlag, Theo Gutman-Solo, Yuhuai Wu, Behnam Neyshabur, Guy Gur-Ari, and Vedant Misra. Solving quantitative reasoning problems with language models, 2022.
- [118] [REDACTED], Jaechan Lee, Yangsibo Huang, Sadhika Malladi, Jieyu Zhao, Ari Holtzman, [REDACTED], [REDACTED], [REDACTED], and Chiyuan Zhang. Muse: Machine unlearning six-way evaluation for language models, 2024.

A Model Details

A.1 Baseline Details

Prompt-based routing We implemented a domain-specific classifier where each input query is categorized into exactly one of n domains using the prompts shown below. Once classified, the query is routed exclusively to the corresponding expert model for processing. In particular, the following prompt was used for our primary experimental setup:

```
"""Classify the given text into one of the following domains:
- 0: Math
- 1: Code
- 2: General science and educational content
- 3: News
- 4: Creative writing and literature
- 5: Reddit
- 6: Research and academic content
- 7: Other

Provide the number between 0 and 7."""
```

For ablations with four datasets instead of eight, we used the following prompting.

```
"""Classify the given text into one of the following domains:
- 0: Math
- 1: Code
- 2: General science and academic content
- 3: Other

Provide the number between 0 and 3."""
```

Branch–Train–Merge (BTM) BTM [11] first assigns a weight to each expert and then combines their predictions through ensembling.

Formally, for a test instance x and the expert i we compute the negative-log-likelihood loss $\mathcal{L}_i(x)$ and convert it to a weight with a temperature-controlled softmax:

$$w_i(x) = \frac{\exp(-\mathcal{L}_i(x)/\tau)}{\sum_{j=1}^M \exp(-\mathcal{L}_j(x)/\tau)},$$

where $\tau > 0$ sharpens ($\tau < 1$) or flattens ($\tau > 1$) the distribution over the M experts. We optionally retain only the k largest-weight experts $S_k(x)$ ($k < M$) and renormalize:

$$\tilde{w}_i(x) = \frac{w_i(x) \mathbf{1}[i \in S_k(x)]}{\sum_{j \in S_k(x)} w_j(x)}.$$

At inference time, let $\mathbf{z}_{i,t} \in \mathbb{R}^{|V|}$ be the logit vector produced by the expert i at generation step t . We combine logits token-by-token:

$$z_t(v) = \sum_{i \in S_k(x)} \tilde{w}_i(x) z_{i,t}(v), \quad v \in V.$$

A.2 Training Details

In FLEXOLMO, to initialize router embeddings with domain embeddings, we sample 1,000 documents from each data source, process them through GritLM/GritLM-7B [13] to obtain document embeddings, and then average these embeddings.

To obtain proxy data $\hat{D}_i$, we train a binary classifier to distinguish D_i from D_{pub} and select public samples with the highest predicted likelihood of belonging to D_i . Specifically, we finetune Snowflake/snowflake-arctic-embed-xs [111], which contains 22M parameters, using a learning rate of 3×10^{-6} . The classifier is trained on a balanced dataset of 500,000 samples (250,000 documents from each source - public and private). The classifier quickly achieved an accuracy above 95% across all datasets considered.

B Data Details

Table 5 presents the statistics of our training data described in Section 4.1.

B.1 Reddit Data Processing

The construction of this dataset involved three major phases.

1. Reddit data filtering A dataset of submission/comment pairs was derived from the Pushshift Reddit dataset [112] (bulk dump as of March 2023) – the same dump used for Dolma Reddit (dolma).

To derive our initial dataset, we extracted each submission and concatenated it with its top-scoring, top-level comment. We then performed further rule-based filtering with the following constraints:

- Filter out deleted/removed content.
- Filter out content marked as over_18.
- Filter out all posts from a list of 26,123 banned or NSFW subreddits.
- Filter out posts from likely bot authors (drawn from <https://botrank.pastimes.eu/> as of Sept 2024).
- Filter out posts containing non-text media.
- Perform document-level text deduplication via Bloom filter.

2. Retrieval-based subreddit selection Dense retrieval was then used to identify academically-relevant subreddits for further filtering. We adapted search queries from MMLU test questions, and performed dense retrieval with these queries on the filtered Reddit data from Step #2, retaining the top 5 hits for each query. Based on these retrieved outputs, we selected 151 subreddits meeting the following criteria:

- Subreddit has ≥ 20 *unique* retrieved items for queries within a given MMLU category; OR
- Subreddit has ≥ 100 retrieved items for queries across all MMLU categories.

We then filtered the dataset from Step #1 to retain only documents from subreddits on this list of 151 subreddits.

3. Format rewriting Finally, the data from Step #2 was input to a synthetic rewriting pipeline to generate academic QA items with coverage of diverse question formats. We defined 7 categories of question format inspired by variation observed in MMLU, and used these to construct prompts for QA text generation. The format categories are as follows:

1. open-ended
2. statement completion
3. fill-in-the-blank
4. statement truth verification
5. which-of-following-has-property-X
6. which-of-following-is-true
7. in-question options

For each format category we constructed a prompt for generating questions of that category given an input text. Below is an example prompt, for the “in-question-options” category. Prompts for other categories differ in 1) the content of the “For format ...” paragraph and 2) the in-context examples (1-3 examples per prompt).

Figure 5: Statistics of our data mix (descriptions in §4.1).

Name	# Tokens (B)
Public Mix	2.37×10^3
News	158.0
Creative Writing	201.9
Math	20.3
StarCoder	83.0
Academic	58.6
Educational text	102.2
Reddit	9.9

I will ask you to convert a text into multiple-choice questions. Here is the text:

“{text}”

Instructions: Convert the information in the text into academic multiple choice questions. ONLY include questions that are academic. DONOT reference the text in the question.

For format, use questions that provide options within the question and give choices for which options are true. Examples:

Dogs have which of the following properties?

- I. They are mammals*
- II. They have five legs.*
- III. They have a tail.*

- A. I only*
- B. II only*
- C. III only*
- D. I and III*

Answer: D

%%%

Which of the following are cities in the US?

- I. Paris*
- II. Athens*
- III. Chicago*

- A. I only*
- B. II only*
- C. III only*
- D. I, II and III*

Answer: C

Separate ALL questions with “\n%%%\n”.

For generating our rewritten QA data, we prompted GPT-4o [REDACTED] (Jan 2025 version). We iterated over the submission/comment pairs in the data from Step #2, and for each of these texts we sampled a format category and prompted the GPT-4o [REDACTED] to generate QA pairs for that text and format category. For longer input texts, format categories were resampled and prompted for again, a number of times proportional to the length of the text.

Finally, GPT-4o [REDACTED] outputs were parsed into separate QA items based on the “%%%

We validated these rewritten data in experiments with OLMo 7B [113] models trained to 2T tokens, carrying out continued pretraining on a 50-50 mix of DCLM and Reddit data while annealing the learning rate to zero, a strategy used in [114, 78, 107]. We run this continued pretraining with two versions of Reddit data: the filtered data from Step #2, and the rewritten data from Step #3. We find that the rewriting improves over the non-rewritten data in both MC9 and MMLU: **MC9 improves from 0.74 to 0.76 and MMLU improves from 0.62 to 0.66.**

Table 4: **Evaluation benchmarks.** . Benchmarks are divided into general-purpose and domain-specific categories. Original dataset citations are listed at right. [†]SciRiFF spans several subtasks (BioASQ factoid, general, and yes/no questions, COVID DeepSet QA, and PubMedQA) evaluated with different metrics. MATH[†] includes seven subsets: algebra, counting and probability, geometry, intermediate algebra, number theory, prealgebra, and precalculus.

General Benchmarks			Domain-specific Benchmarks		
Benchmark	Metric	Citation	Benchmark	Metric	Citation
ARC-EASY	Acc _{raw}	[84]	HUMANEVAL	Pass@1	[105]
ARC-CHALLENGE	Acc _{raw}	[84]	HUMANEVALPLUS	Pass@1	[104]
BOOLQ	Acc _{raw}	[85]	MBPP	Pass@1	[103]
CSQA	Acc _{raw}	[86]	MBPPPLUS	Pass@1	[104]
HELLASWAG	Acc _{raw}	[87]	SciRiFF [†]	—	[106]
OPENBOOKQA	Acc _{raw}	[88]	MATH [†]	Exact Match	[102]
PIQA	Acc _{raw}	[89]	NEWS GENERATION	LM-as-judge	-
SOCIAL IQA	Acc _{raw}	[90]	POEM GENERATION	LM-as-judge	-
WINOGRANDE	Acc _{raw}	[91]	GSM8K	Exact Match	[101]
MMLU	Acc _{raw}	[97]			
MMLU-PRO	Acc _{raw}	[98]			
AGIEVAL (English)	Acc _{raw}	[99]			
BIG-BENCH HARD (BBH)	Exact Match	[100]			
CoQA	F1	[92]			
DROP	F1	[96]			
NATURAL QUESTIONS (NQ)	F1	[94]			
SQUAD	F1	[93]			
TRIVIAQA	F1	[95]			
NARRATIVEQA	F1	[115]			

Aggregate scores (averages of individual benchmarks)
GEN5: CoQA, SQUAD, NATURAL QUESTIONS, TRIVIAQA, DROP
MC9: ARC-EASY, ARC-CHALLENGE, BOOLQ, CSQA, HELLASWAG, OPENBOOKQA, PIQA, SOCIAL IQA, WINOGRANDE
CODE4: MBPP, MBPPPLUS, HUMANEVAL, HUMANEVALPLUS

C Evaluation Details

All evaluations are done using the OLMES evaluation standard introduced by [116], following key metrics from the OLMo 2 framework [78]. Table 4 presents a detailed breakdown of our evaluation datasets and metrics. General description is provided in §4.2; here, we provide more details.

SciRiFF We select five subtasks that do not require structured prediction: BioASQ-Factoid, BioASQ-Yes/No, BioASQ-General, PubMedQA, and COVID-QA. , with our tables reporting the average performance across all five tasks.

MATH We assess solution correctness through exact matching with ground truth answers, following the methodology established in [117].

News Generation We use the `muse-bench/MUSE-News` dataset from [118], selecting articles containing between 64 and 128 tokens. Models are prompted to continue an article given a prefix of the first 32 tokens, using a sampling temperature of 0.8. A `Llama-3.3-70B-Instruct` model serves as the judge, evaluating each completion based on journalistic quality (2 points), topical coherence (2 points), and clarity/fluency (1 point), for a total score between 0 and 5, which is then normalized to a 0–100 scale. To reduce variance, we generate five completions per prompt and report the average score.

Poem Generation We employ the `merve/poetry` dataset⁵, filtering poems to those containing 64–128 tokens. From a total of 176 poems (147 Renaissance, 29 Modern), we reserve five poems (three Renaissance, two Modern) for few-shot examples and use 100 for testing. Models continue each poem from its first four lines, following the same prompting and evaluation settings as in NEWSGEN. A genre-aware `Llama-3.3-70B-Instruct` judge evaluates each completion based on poetic craftsmanship (2 points), thematic coherence (2 points), and clarity/fluency (1 point), with

⁵<https://www.kaggle.com/datasets/ishnoor/poetry-analysis-with-machine-learning/data>

scores normalized to a 0–100 scale. As with news generation, five completions are generated per prompt, and we report the average score across all instances.

D Methodology Intuition

D.1 Problem Setup

Given an input data $\mathbf{x} \in \mathbb{R}^h$, the router learning can be viewed as a multi-class classification problem with $n+1$ classes: a public class $\mathcal{C}_{\text{pub}}$ and n closed classes $\{\mathcal{C}_i\}_{i=1}^n$. The scoring function $s_i : \mathbb{R}^h \rightarrow \mathbb{R}$ such that $s_i(\mathbf{x}) = \mathbf{r}_i \cdot \mathbf{x}$ represents the score for class $\mathcal{C}_i$, where $\mathbf{r}_i$ is the router embedding and the higher scores indicate higher class membership likelihood.

Training: Pairwise Binary Classification In training experts to coordinate (§3.3.1), we learn n binary classifiers $\{f_i\}_{i=1}^n$, where each $f_i : \mathbb{R}^h \rightarrow \{\mathcal{C}_{\text{pub}}, \mathcal{C}_i\}$ discriminates between $\mathcal{C}_{\text{pub}}$ and $\mathcal{C}_i$:

$$f_i(\mathbf{x}) = \begin{cases} \mathcal{C}_i & \text{if } s_i(\mathbf{x}) > s_{\text{pub}}(\mathbf{x}) \\ \mathcal{C}_{\text{pub}} & \text{otherwise.} \end{cases}$$

The decision boundary h_i between $\mathcal{C}_{\text{pub}}$ and $\mathcal{C}_i$ is defined as:

$$h_i := \{\mathbf{x} \in \mathbb{R}^h : s_{\text{pub}}(\mathbf{x}) = s_i(\mathbf{x})\}.$$

Inference: Multiclass Classification At inference time (§ 3.3.2), these n binary classifiers are combined into a unified multi-class classifier that maps from $\mathbb{R}^h$ to the complete class space $\mathcal{C}_{\text{pub}}, \mathcal{C}_1, \dots, \mathcal{C}_n$. The final classification decision is determined by:

$$F(\mathbf{x}) = \arg \max_{i \in \{\text{pub}, 1, \dots, n\}} s_i(\mathbf{x}).$$

The key question is: how could we make the ensemble of binary classifiers $\{f_i\}_{i=1}^n$ route inputs as close to if we had trained one unified multiclass classifier F end-to-end on all data?

Our intuition is if each binary classifier f_i learns the decision boundary " $\mathcal{C}_i$ vs. not $\mathcal{C}_i$ " rather than just " $\mathcal{C}_i$ vs. $\mathcal{C}_{\text{pub}}$ ", this could make decision boundary that tightly encircles $\mathcal{C}_i$ which leads to better multi-class classification.

D.2 Anchor Point and Negative Bias

To ensure each classifier f_i learns a better decision boundary h_i , we propose two key techniques:

Freezing $\mathbf{r}_{\text{pub}}$ and M_{pub} as anchor points During training of each binary classifier f_i , we fix the public expert M_{pub} and router embedding $\mathbf{r}_{\text{pub}}$ to ensure all classifiers share a common coordinate system by maintaining:

$$s_{\text{pub}}(\mathbf{x}) = \mathbf{r}_{\text{pub}} \cdot \mathbf{x}, \forall i \in \{1, \dots, n\}.$$

Without this constraint, each binary classifier would optimize both $\mathbf{r}_{\text{pub}}$ and $\mathbf{r}_i$ independently, resulting in inconsistent coordinate systems where $\mathbf{r}_{\text{pub}}^{(i)} \neq \mathbf{r}_{\text{pub}}^{(j)}$ for $i \neq j$ during the training of binary classifiers f_i and f_j . Such inconsistency would invalidate the multi-class classifier $F(\mathbf{x})$, as it would attempt to compare scores computed in incompatible embedding spaces. By maintaining a fixed reference point, we ensure all decision boundaries h_i are defined relative to the same coordinate system, enabling their meaningful composition during inference without additional training.




Figure 6: During training, the negative bias shifts the decision boundary. So that a more selective subset of data will be used to train the expert corresponding to $\mathcal{C}_1$.

Adding a negative bias During training of f_i , each binary classifier only needs to satisfy $\mathbf{r}_i \cdot \mathbf{x} > \mathbf{r}_{\text{pub}} \cdot \mathbf{x}$ for $\mathbf{x} \in \mathcal{C}_i$. When $\mathcal{C}_{\text{pub}}$ and $\mathcal{C}_i$ is separable, many possible $\mathbf{r}_i$ vectors (no bias boundary in Figure 6) can satisfy this constraint without precisely characterizing the specified region of $\mathcal{C}_i$. We refine this by adding a negative bias term to each expert:

$$\mathbf{r}_i \cdot \mathbf{x} + b_i > \mathbf{r}_{\text{pub}} \cdot \mathbf{x} \quad \forall i \in \{1, 2, \dots, n\}$$

As Figure 6 illustrates, the negative bias term could help move the decision boundary h_i closer to $\mathcal{C}_i$'s data points. This means, during training, a more selective subset of data will be used to train the local expert. This facilitates merging, where experts compete not only with $\mathcal{C}_{\text{pub}}$ but with each other.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction accurately describe our core contributions.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Please refer to the result analysis §5.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[NA\]](#)

Justification: No formal results.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: The code will be open-sourced and we describe all model experiments in detail.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Data and code will be open-sourced with instructions

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: See hyperparameter in §A

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: Repeated experimental runs to produce error bars would be prohibitively expensive; we instead capture evaluation noise through inclusion of a large suite of downstream benchmark scores, see §4.2.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).

- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Yes, we report computer resources in detail in §4.4.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [No] con

Justification: The research uses restricted datasets specifically to demonstrate the data flexibility and privacy compliance of the proposed model architecture FLEXOLMO .

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: We discuss these in §2.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.

- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [\[Yes\]](#)

Justification: Yes, the paper describes safeguards that allow data opt-out at inference time.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: Yes, we describe licenses in §4.1.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.

- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: Yes, documentation is provided throughout the text.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[NA\]](#)

Justification: There are no crowdsourcing experiments nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the ████ minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[NA\]](#)

Justification: There are no crowdsourcing experiments nor research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]

Justification: The core method development in this research does not involve LLMs in non-standard components.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.